

ARTICLE

Learnable Role-Conditioned Evidence Sentence Selection with Interpretable Mixture Reranking

Bijing Liu^{1,2} and Yong Yang^{1,2,*}

¹NARI Group Corporation (State Grid Electric Power Research Institute), Nanjing 211106, Jiangsu Province, China

²Beijing Kedong Electric Power Control System Co., Ltd., Beijing 100080, China

*Corresponding Author: Yong Yang. Email: yangyong1@sgepri.sgcc.com.cn

Version August 4, 2026 submitted to Comput Mater Contin

ABSTRACT: Evidence sentence selection asks a model to return the supporting sentence IDs for a claim or role-conditioned question, which is a core building block for auditable retrieval and fact verification. This paper introduces C²GES, an interpretable mixture reranker that combines frozen query relevance, a learnable role-compatibility head, and a light local-chain consistency signal. Unlike lexicon-only role scoring, the role head is trained on held-out human-gold labels with a pairwise ranking objective and mixture-weight floors that prevent the role term from collapsing. The main quantitative evaluation uses the filtered FEVER evidence-selection release with human-annotated evidence lines: 4,000 training, 800 development, and 800 test claim–document instances under SUPPORTS/REFUTES role conditioning. On the test split with budget $K = 3$, C²GES reaches 0.5066 evidence F1, improving over TF-IDF (0.4837), SBERT (0.4818), a lexicon-cue role baseline (0.4414), and the query-only ablation (0.4937). Document-cluster bootstrap shows a significant gain over the no-role ablation (+0.0103, 95% CI [0.0027, 0.0187], $p = 0.005$) and over TF-IDF/SBERT ($p < 0.001$), while remaining statistically tied with BM25. Public North American Electric Reliability Corporation (NERC) reliability-report excerpts are used only as qualitative case illustrations of role-conditioned retrieval in power-grid technical prose, not as gold-label leaderboards. The contribution is therefore a learnable yet auditable role-conditioned reranker validated on a public human-gold evidence benchmark, with domain transfer demonstrated through NERC case inspection.

KEYWORDS: Evidence retrieval; fact verification; FEVER; role-conditioned reranking; learnable ranking; power grid reliability reports

1 Introduction

Many technical workflows require not only a retrieved answer, but the exact supporting sentences that justify it. Fact verification, query-focused summarization, and post-event report review all benefit from sentence-ID evidence selection because the selected spans remain inspectable and reusable. FEVER established a large-scale human-annotated setting in which claims are linked to Wikipedia evidence sentences [1]. Parallel work in query-focused and extractive summarization shows that conditioning on an explicit information need changes which sentences should be returned [2–4]. In power-system reliability analysis, the same requirement appears when engineers inspect disturbance and lessons-learned reports: a trigger, impact, or mitigation question may share topical vocabulary while needing different evidence sentences [5,6].

Generic dense retrieval and lexicon heuristics leave a methodological gap. Lexical cue lists are easy to interpret but hard to justify when they are handcrafted on the evaluation domain. Pure dense matching is

stronger on average but does not expose an explicit role or claim-type factor that can be ablated. C²GES addresses this gap with an additive, auditable score

$$S(s, q, r) = w_q Q(s, q) + w_r R_\theta(s, q, r) + w_g G(s, D, r), \quad (1)$$

where Q is frozen query relevance, R_θ is a learnable role-compatibility head, and G is a light local-chain consistency feature. Mixture weights are constrained with floors so that the role term cannot be driven to zero during training.

This paper makes three contributions:

- It introduces C²GES as a role-conditioned evidence reranker with a learnable compatibility head trained under held-out human-gold supervision, retaining an interpretable linear mixture for ablation.
- It evaluates the method on the filtered FEVER evidence-selection benchmark with human-annotated evidence sentence IDs, reporting baselines, ablations, and document-cluster bootstrap tests.
- It illustrates domain transfer on public NERC reliability-report excerpts as qualitative cases, without treating those excerpts as a gold-label leaderboard.

The remainder of the paper is organized as follows. Section 2 reviews related work. Section 3 formalizes the task and FEVER conversion. Section 4 presents the learnable C²GES model. Section 5 reports experiments and NERC case illustrations. Section 6 concludes.

2 Related Work

2.1 Evidence Selection and Fact Verification

FEVER and its follow-on shared tasks treat evidence retrieval as a first-class objective alongside claim labeling [1]. Subsequent reranking and retrieval studies show that sentence-level evidence recovery remains challenging even when claim classification is strong. This paper uses a filtered FEVER evidence-selection release in which each instance provides a document sentence pool and human gold evidence line IDs, converted to the same top- K sentence-ID metric used by extractive evidence systems.

2.2 Query-Focused and Extractive Ranking

Query-focused summarization conditions content selection on an information need [2,7,8]. Extractive methods based on matching, centrality, and rank fusion provide transparent sentence selection [3,9,10]. C²GES follows the extractive tradition but adds an explicit role or claim-type factor that is learned rather than hand-listed.

2.3 Graph and Chain Consistency Signals

Heterogeneous and multiplex graph summarizers motivate structural sentence interactions [11,12]. C²GES does not train a graph neural network; it uses a proximity-smoothed consistency feature as a small auxiliary term so that ablations isolate role learning from local structure.

2.4 Power-Grid Text Analytics

AI methods for grid modeling, reliability, and digital infrastructure motivate domain transfer to NERC-style reports [5,6,13,14]. In this paper, NERC text is used for qualitative inspection only; quantitative claims rely on FEVER human-gold evidence.

3 Task and Benchmark 66

3.1 Sentence-ID Evidence Selection 67

Let a document be a sentence sequence $D = \{s_1, \dots, s_n\}$. Given a query q and a discrete role r , the system returns a set of at most K sentence IDs $\hat{Y}^K$. Evaluation compares $\hat{Y}^K$ with a gold evidence set Y by precision, recall, and F1 over sentence IDs. The primary operating point is $K = 3$. 68
69
70

3.2 FEVER Human-Gold Conversion 71

The main benchmark is built from the filtered FEVER evidence-selection dataset. Each instance contains a claim, a SUPPORTS or REFUTES label, a Wikipedia document sentence pool, and human-annotated evidence line IDs. We map SUPPORTS/REFUTES to roles $\{supports, refutes\}$, convert wiki line IDs to sentence IDs $s_{000}, s_{001}, \dots$, and treat the claim as the role-conditioned question. After conversion, the experimental splits contain 4,000 training, 800 development, and 800 test instances. Labels are human gold from FEVER; no model-generated labels are used for training or testing. 72
73
74
75
76
77

3.3 NERC Case Material 78

Public NERC disturbance and lessons-learned report excerpts are retained only for qualitative case illustrations of how role-conditioned ranking behaves on power-grid technical prose. Those excerpts are not used as a gold-label test set in the reported tables. 79
80
81

4 Method 82

4.1 Interpretable Mixture 83

For each sentence s_j , C^2 GES computes 84

$$S_j = w_q Q_j + w_r R_j + w_g G_j, \quad (2) \quad 85$$

then returns the top- K sentence IDs. Mixture weights (w_q, w_r, w_g) are positive and constrained by floors $(0.35, 0.25, 0.05)$ before renormalization, which keeps the role term active during learning. 86
87

4.2 Frozen Query Relevance 88

Q_j combines SBERT cosine similarity between sentence and claim embeddings with TF-IDF cosine similarity. Both channels are min-max normalized within the candidate pool and averaged. This keeps the backbone transparent and avoids co-adapting a deep retriever with the role head. 89
90
91

4.3 Learnable Role Compatibility 92

Let h_s and h_q be frozen MiniLM embeddings and let e_r be a learned role embedding. The role head is an MLP 93
94

$$R_j = \sigma(\text{MLP}([h_{s_j}; h_q; e_r])). \quad (3) \quad 95$$

Unlike a manual cue lexicon, R_θ is trained on held-out FEVER gold evidence with a pairwise ranking loss. A lexicon-cue scorer is retained only as a non-learned baseline. 96
97

4.4 Local-Chain Consistency

G_j is a proximity-smoothed salience feature over neighboring sentences. It is intentionally weak; the no-graph ablation isolates whether local structure adds value beyond query and role scores.

4.5 Training Objective

Training maximizes the score margin between gold and non-gold sentences under the full mixture, plus an auxiliary pairwise loss on the role head alone. Development F1 selects the checkpoint. Document instances used for testing are held out from training.

Require: Sentences D , claim q , role r , budget K , trained head/weights

- 1: **for** each sentence $s_j \in D$ **do**
- 2: Compute Q_j, R_j, G_j
- 3: $S_j \leftarrow w_q Q_j + w_r R_j + w_g G_j$
- 4: **end for**
- 5: Sort by S_j descending
- 6: **return** top- K sentence IDs

5 Experiments and Discussion

5.1 Setup

All methods share the same FEVER-converted sentence pools and $K = 3$ budget. Baselines are lead- K , TF-IDF, BM25, SBERT, and a lexicon-cue role scorer. C²GES ablations are query-only, no-role, and no-graph. Uncertainty uses document-cluster bootstrap over test instances (400 resamples). The selected mixture on the best development checkpoint is approximately (0.70, 0.25, 0.05).

5.2 Main Results

Table 1 reports test evidence F1. Full C²GES attains 0.5066 F1, ahead of TF-IDF (0.4837), SBERT (0.4818), lead- K (0.4590), and lexicon cues (0.4414). BM25 is close at 0.5030. The learned role head is the main methodological gain: removing it drops F1 to 0.4967, and the bootstrap difference is significant ($p = 0.005$). Gains over TF-IDF, SBERT, query-only, and lexicon cues are also significant ($p < 0.001$). The BM25 comparison is not significant.

Table 1: FEVER evidence-selection test results ($K = 3$; human-gold evidence; best F1 in bold)

Method	Precision	Recall	F1
Lead- K	0.3594	0.7754	0.4590
LexCue role	0.3398	0.7577	0.4414
TF-IDF	0.3740	0.8293	0.4837
SBERT	0.3698	0.8291	0.4818
BM25	0.3881	0.8631	0.5030
C ² GES query-only	0.3810	0.8478	0.4937
C ² GES no-role	0.3840	0.8508	0.4967
C ² GES no-graph	0.3898	0.8636	0.5038
C²GES full	0.3919	0.8680	0.5066

Table 2: Document-cluster bootstrap comparisons of full C²GES versus selected alternatives on FEVER test

Comparison	Mean Δ F1	95% CI	p
vs TF-IDF	+0.0234	[0.0137, 0.0336]	< 0.001
vs SBERT	+0.0250	[0.0150, 0.0346]	< 0.001
vs query-only	+0.0133	[0.0055, 0.0216]	< 0.001
vs no-role	+0.0103	[0.0027, 0.0187]	0.005
vs LexCue	+0.0653	[0.0522, 0.0790]	< 0.001
vs BM25	+0.0041	[-0.0049, 0.0128]	0.365

Role-stratified F1 is 0.5167 on SUPPORTS and 0.4977 on REFUTES, indicating that the learned head helps both claim types rather than a single label.

5.3 Ablations and Interpretation

The no-role ablation is the decisive integrity check for the proposed contribution. Because the role head is trained on held-out human gold and remains active under mixture floors, its removal hurts evidence F1 significantly. The graph term contributes a smaller point-estimate gain (full 0.5066 vs no-graph 0.5038). LexCue underperforms because fixed surface cues do not transfer cleanly across Wikipedia domains, which motivates replacing handcrafted dictionaries with a learned role embedding.

5.4 NERC Qualitative Cases

On public NERC solar-PV disturbance and winter-storm report excerpts, role-conditioned ranking tends to surface initiating fault sentences for trigger-like queries and recommendation or restoration language for mitigation-like queries, even when both mention outages or MW impacts. These observations are illustrative only: without an independent human-gold NERC evidence set, we do not report NERC F1 as a primary result.

5.5 Discussion

The results support a scoped claim: a learnable role-compatibility head inside an interpretable mixture improves FEVER evidence-sentence recovery over query-only and lexicon-role baselines, while remaining competitive with strong lexical retrieval. The method is intentionally lighter than end-to-end neural rerankers; its value is auditability and an isolatable role factor. Absolute F1 remains far from saturation, so denser candidates, cross-encoders, and multi-hop evidence sets remain important future baselines.

6 Conclusions, Limitations, and Future Work

C²GES combines frozen query relevance, a learnable role head, and a light chain feature into an auditable evidence-sentence reranker. On human-gold FEVER evidence selection it improves over query-only and lexicon-role baselines, with a significant no-role ablation gap, and stays competitive with BM25. NERC report excerpts demonstrate qualitative domain transfer without serving as gold-label benchmarks.

Limitations include the use of a filtered single-document FEVER setting rather than full open retrieval over the entire Wikipedia dump, a moderate training subset (4,000 instances), frozen rather than jointly trained encoders, and the absence of a human-gold power-grid evidence corpus. Future work should add larger FEVER open-retrieval settings, cross-encoder baselines, span-aware metrics, and carefully collected expert evidence labels for reliability reports.

Acknowledgement: During the preparation of this work, the authors used large-language-model-based tools to assist with drafting and editing. After using these tools, the authors reviewed and edited the content as needed and take full responsibility for the content of the published article.

Funding Statement: This work was supported by the Science and Technology Project of NARI Group Corporation (State Grid Electric Power Research Institute) (Grant No. [AUTHOR INPUT REQUIRED: grant number]).

Author Contributions: The authors confirm contribution to the paper as follows: Conceptualization, methodology, software, validation, data curation, and writing—original draft preparation: L.B.; supervision, project administration, resources, funding acquisition, and writing—review and editing: Y.Y. All authors reviewed and approved the final version of the manuscript.

Availability of Data and Materials: The FEVER-derived sentence-ID benchmark conversion scripts, trained checkpoints, and evaluation summaries are available in the project workspace under `workspace/fever_benchmark/` and `workspace/fever_runs/learnable_role/`. Public FEVER resources are available at <https://fever.ai/dataset/fever.html>. Public NERC reports used for qualitative cases are available from NERC event-analysis pages.

Ethics Approval: Not applicable. This study did not involve human or animal subjects; FEVER labels are from a public human-annotated benchmark, and NERC materials are public technical reports.

Conflicts of Interest: The authors declare no conflicts of interest.

Supplementary Materials: The supplementary material is available online at <https://www.techscience.com/doi/10.32604/cmc.2026.000000/s1>. It documents data conversion details, hyperparameters, and additional split statistics.

References

- Thorne J, Vlachos A, Christodoulopoulos C, Mittal A. FEVER: a large-scale dataset for fact extraction and VERification. In: Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers). New Orleans, Louisiana; 2018. [CrossRef].
- Zhong M, Yin D, Yu T, Zaidi A, Mutuma M, Jha R, et al. QMSum: a new benchmark for query-based multi-domain meeting summarization. In: Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies; 2021. [CrossRef].
- Zhong M, Liu P, Chen Y, Wang D, Qiu X, Huang X. Extractive summarization as text matching. In: Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics; 2020. [CrossRef].
- Vig J, Fabbri AR, Kryściński W, Wu CS, Liu W. Exploring neural models for query-focused summarization. In: Findings of the Association for Computational Linguistics: NAACL 2022; 2022. [CrossRef].
- Xie L, Zheng X, Sun Y, Huang T, Bruton T. Massively digitized power grid: opportunities and challenges of use-inspired AI. Proc IEEE. 2023;111(7):762-87. [CrossRef].
- Madabhushi S, Dewri R. A survey of anomaly detection methods for power grids. Int J Inf Secur. 2023;22(6):1799-832. [CrossRef].
- Liu X, Xu Y. Learning to rank utterances for query-focused meeting summarization. arXiv:2305.12753. 2023.
- Sotudeh S, Goharian N. Learning to rank salient content for query-focused summarization. arXiv:2411.00324. 2024.
- Bi K, Jha R, Croft WB, Çelikylmaz A. AREDSUM: adaptive redundancy-aware iterative sentence ranking for extractive document summarization. In: Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics; 2021. [CrossRef].
- Khor YK, Tan CW, Lim TM. Text summarization on Amazon food reviews using TextRank. In: Proceedings of the International Conference on Digital Transformation and Applications (ICDXA); 2021. [CrossRef].

11. Wang D, Liu P, Zheng Y, Qiu X, Huang X. Heterogeneous graph neural networks for extractive document summarization. In: Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics; 2020. [[CrossRef](#)]. 197
198
199
12. Jing B, You Z, Yang T, Fan W, Tong H. Multiplex graph neural network for extractive text summarization. In: Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing; 2021. [[CrossRef](#)]. 200
201
13. Srinivasan S, Kumarasamy S, Andreadakis Z, Lind PG. Artificial intelligence and mathematical models of power grids driven by renewable energy sources: a survey. *Energies*. 2023;16(14):5383. [[CrossRef](#)]. 202
203
14. Hamann HF, Brunswiler T, Gjorgiev B, Martins LSA, Puech A, Varbella A, et al. Foundation models for the electric power grid. *arXiv:2407.09434*. 2024. 204
205