

# C<sup>2</sup>GES: A Diagnostic Evaluation of Role-Conditioned Extractive Summarization for Long Power-System Technical Reports

Bijing Liu <sup>1,2</sup> and Yong Yang <sup>1,2,\*</sup>

<sup>1</sup> NARI Group Corporation (State Grid Electric Power Research Institute), Nanjing 211106, Jiangsu Province, China;

<sup>2</sup> Beijing Kedong Electric Power Control System Co., Ltd., Beijing 100080, China

\* Correspondence: yangyong1@sgepri.sgcc.com.cn

## Abstract

Structural additions to an extractive selector can increase complexity without improving quality. As an information question: does added structure supply what lexical relevance and dense similarity do not already carry? We present C<sup>2</sup>GES, a deterministic selector with lexical relevance, role heuristics, a typed proxy graph, path participation, redundancy control, and page traceability. On 15 NERC reports, Full C<sup>2</sup>GES had higher ROUGE-L than Semantic-MMR and TextRank at equal unit counts, but outputs were 54–63% longer and no matched-word contrast survived Holm adjustment. A post-access seven-series pilot at 110/260-word budgets placed TextRank first at 110 words (0.1690) and no-path C<sup>2</sup>GES first at 260 words (0.1844); Full-minus-no-path was -0.00799/-0.00718. A freeze-before-eval path redesign scored 0.0652/0.1044 on the same 15 reports at those budgets against no-path 0.0654/0.1020 and TextRank 0.0675/0.1031; this is held-out for this revision, not unseen confirmatory, and is not a win on both budgets. On eight fictional role-chain fixtures, Full scored 0.3886/0.3481 versus no-path 0.3750/0.3269 (Holm 1.0); those fixtures are not real incident data. The evidence supports a negative real-summary diagnosis and no-path as the simpler provisional default, but not system superiority, structural-proxy validity, or operational benefit.

**Keywords:** extractive summarization; diagnostic evaluation; negative results; power-grid reports; length-controlled evaluation; reproducibility

## 1. Introduction

Power-system reliability and event-analysis reports often extend across hundreds of pages. Evidence about an initiating condition, a trigger, system propagation, consequences, and corrective action may appear in different sections and coexist with tables, standards language, recommendations, and appendices. Engineers reviewing such material need more than a list of individually prominent sentences. They need a compact route through structurally complementary passages and a reliable way to return to the source.

Extractive summarization is attractive in this setting because it preserves the wording and location of selected passages. Position, centroid-style lexical relevance, graph centrality, and semantic diversity provide useful ranking signals [1–5]. These signals, however, do not necessarily distinguish the roles that passages play in an engineering account. A highly central passage may repeat background information, whereas a less central passage may contain the only mitigation or consequence needed to interpret an event.

Graph-based summarizers model relationships among sentences through learned graphs, hypergraphs, or structural filtering [6–8]. Power-domain language studies use

Received:

Accepted:

Published:

**Copyright:** © 2026 by the authors.

Submitted to *Information* for possible open access publication under the terms and conditions of the [Creative Commons Attribution \(CC BY\)](#) license.

entity–relation extraction and knowledge graphs to organize fault, alarm, and inspection evidence [9–11]. C<sup>2</sup>GES links these lines of work in an engineering-oriented formulation: candidate units receive transparent textual roles, compatible roles form typed paths, and path participation is combined with relevance, graph salience, position, and redundancy control. Each selected unit remains linked to its original report page.

The study uses complete public reliability, disturbance, event-analysis, recommendation, and assessment reports from the North American Electric Reliability Corporation (NERC). NERC guidance on the attributes of quality event-analysis reports provides the institutional context for this report family [12]. These reports contain event narratives and corrective lessons but are not utility work orders, inspection records, or field-maintenance narratives. The corpus is therefore treated as a maintenance-oriented technical-report proxy. Twelve reports support development, and 15 reports form the retained test set. All candidates are drawn from report pages outside the official Executive Summary, which is used only as the reference for automatic evaluation.

Beyond its summarization framing, the question this study turns on is an information one: does a typed structural representation supply discriminative information that lexical relevance and dense embedding similarity do not already carry, or does it merely re-encode what those signals already supply? The distinction has consequences beyond this corpus: if the typing is redundant, the graph machinery adds representational cost without adding extractable information, and a simpler untyped representation is the better instrument. This diagnostic study asks three bounded questions. *RQ1* asks whether the apparent retained-test advantage persists after output length and report-series dependence are examined and whether the same ordering appears in a separately identified post-access pilot with matched 110- and 260-word budgets. *RQ2* asks which role, reservation, graph, path, and redundancy components change performance or selection behavior in controlled exploratory variants. *RQ3* asks what extraction defects and role disagreements are exposed by an explicitly non-human two-model audit, without treating machine agreement as construct validation.

The paper makes three current contributions. First, it provides a deterministic, source-linked role-conditioned selector whose scoring channels can be inspected and removed under controlled configurations. Second, it reports a layered evaluation that separates the 15-report retained test from a seven-series post-access pilot and explicitly diagnoses length, clustering, layout, embedding, and multiplicity effects. Third, it reports a reproducible negative component result: path deletion did not improve the predefined retained-test endpoint, was negative relative to no-path in the exploratory word-budget pilot, and typed edge information showed no stable exploratory advantage over an untyped graph. Taken together, these results do not establish that role typing contributes extractable information beyond the untyped graph; on the evidence collected, the simpler representation is the defensible default. Stronger confirmatory and semantic-validity claims are reserved for future independently collected evidence and are not contributions of this submission.

## 2. Related Work

### 2.1. Extractive Technical-Document Summarization

Extractive methods rank source sentences using position, document-centroid similarity, graph centrality, learned salience, or redundancy-aware relevance [1–3]. Extractive summarization is itself an information-extraction process, and the question this study measures is which of those signals contribute independently to the measured endpoint. Lead and centroid methods remain useful controls, while TextRank and LexRank provide widely used graph-ranking references [4,13]. Dense sentence representations extend lexical matching and support semantic relevance–diversity selection [5]. Holding candidates, references, preprocessing, and extraction-unit budgets constant allows these ranking principles to be

compared within one report corpus, although representation and tuning opportunities still differ.

PacSum augments thresholded sentence centrality with asymmetric weights for preceding and following positions [14]. It is included as a modern unsupervised comparator because it combines semantic similarity, graph centrality, and position without supervised summary labels. The controlled implementation is explicitly named PacSum-MiniLM: it independently implements the published centrality equations while using the same fixed MiniLM representation and long-unit policy as Semantic-MMR in the exploratory pilot. It is therefore not presented as a byte-equivalent reproduction of the authors' task-fine-tuned BERT implementation.

Long technical documents add layout and discourse problems that are less prominent in short news articles. Bug-report summarization, for example, must distinguish problem statements, observed behavior, diagnostic discussion, and resolution-oriented passages [15]. Multi-document methods must likewise balance common content against source-specific evidence [16]. NERC reports combine long narrative sections with tables, standards language, recommendations, and appendices. This motivates a selector that considers both sentence relevance and the role of a passage in the report's engineering account.

## 2.2. *Graphs, Semantic Ranking, and Causal Language*

Graph summarizers model relations among sentences, entities, topics, or discourse units [17–19]. Learned sentence graphs, contextual hypergraphs, independent-set filtering, and heterogeneous long-document graphs all extend isolated relevance scoring [6–8,20]. StrucSum injects sentence-graph structure into prompts for zero-shot extractive summarization of long documents and reports that combining its structural prompting strategies does not clearly improve performance [21]. That null concerns prompt-level combinations on general-domain corpora rather than typed against untyped edges on a domain report family, but it is a negative structural result in the same problem setting, and we record it as the closest published analogue of the component question asked here. C<sup>2</sup>GES takes a complementary approach. It uses fixed lexical roles and deterministic typed-path enumeration so that the contribution of each score channel and selected sentence can be inspected.

The graph is rebuilt for each report and does not store a persistent asset ontology. Its node-deletion term measures how much qualified textual-path utility is lost when one sentence node is removed. This is a structural sensitivity feature, not an alternative physical event. Causal NLP distinguishes language association from intervention and counterfactual identification [22,23]; accordingly, the present method is described as causal-chain-inspired only in the sense of its condition–event–impact–mitigation role progression. The remainder of the paper uses the more precise term *path-deletion term*.

## 2.3. *Power-Grid Text Analytics*

Grid language technologies cover infrastructure analysis, digital twins, supervised entity–relation extraction, and persistent knowledge graphs [24–29]. Power-fault, line-loss, and alarm-tracing graphs organize entities or events [9,10,30]. Maintenance-oriented retrieval and knowledge-graph studies use equipment or inspection records and sometimes include annotation or practitioner evaluation [11,31,32]. Relation-extraction methods similarly predict supervised entity–relation labels [33,34].

C<sup>2</sup>GES addresses a different level of analysis: it ranks passages from a complete report while preserving source locations and encouraging complementary textual roles. Table 1 positions this formulation against selected adjacent studies. The contribution evaluated here is the integration of deterministic role evidence, typed-path participation, source-linked extraction, and a strict component ablation. The study claims neither absolute

priority over adjacent methods nor a performance benefit from the current path-deletion term.

**Table 1.** Positioning against selected studies with related structural or power-domain objectives.

| Study family                           | Primary object                   | Structural mechanism                             | Evaluation emphasis                                                 | Difference from this study                                                                   |
|----------------------------------------|----------------------------------|--------------------------------------------------|---------------------------------------------------------------------|----------------------------------------------------------------------------------------------|
| Sentence Graph Attention / MCHES [6,7] | General-domain documents         | Learned sentence graph or contextual hypergraph  | Summary metrics on benchmark corpora                                | Learned semantic representations; no NERC proxy or node-deletion ablation                    |
| Graph MIS + LLM [8]                    | DUC CNN/DailyMail and documents  | Independent-set filtering and graph metrics      | ROUGE and computational efficiency                                  | Different graph objective, corpora, and generation components                                |
| Bug-report summarization [15]          | Software issue reports           | Learned sentence significance                    | Technical-artifact summary quality                                  | Similar long technical discourse, but different roles, references, and domain                |
| Power fault/alarm graphs [9,10]        | Power-domain entities and events | Ontology/KG extraction and graph tracing         | Entity, relation, or tracing performance                            | Structured knowledge construction rather than report-level extractive summarization          |
| Maintenance KG and retrieval [11,31]   | Equipment or inspection records  | LLM retrieval or ontology-constrained extraction | Retrieval, annotation, manual audit, or practitioner utility        | Title-concordant data and human assessment; no comparable extractive ablation                |
| C <sup>2</sup> GES                     | Selected NERC technical reports  | Deterministic role graph and typed-path deletion | Executive-Summary overlap, paired contrasts, and length diagnostics | Report-level extractive selection with source-linked outputs and a strict component ablation |


3. Materials and Methods

3.1. Task Formulation and System Overview

Let report  $D = \{s_1, \dots, s_n\}$  contain non-summary candidate extraction units. For budget  $K \in \{5, 10\}$ , method  $m$  returns  $\hat{Y}_{m,K} \subseteq D$ , with  $|\hat{Y}_{m,K}| = \min(K, n)$ . The selected units are restored to source order. The official Executive Summary is used only as the reference and is never candidate input.

ROUGE-1, ROUGE-2, and ROUGE-L F1 quantify lexical overlap between the extract and that Executive Summary [35]. The acronym ROUGE denotes Recall-Oriented Understudy for Gisting Evaluation. ROUGE-L is the predefined primary endpoint; redundancy is a descriptive secondary outcome. These measures characterize reference overlap and lexical repetition rather than expert-assessed completeness or operational usefulness.

Figure 1 summarizes the deterministic pipeline. Complete reports first pass candidate and reference-separation gates. Each retained candidate receives five lexical role scores and contributes, when a unique role is assigned, to a typed sentence graph. Relevance, role, graph, path-deletion, and position channels form a base score. The selector reserves available evidence from three role groups before filling the remaining budget with a redundancy-aware greedy rule. Selected passages are restored to source order and retain identifiers that resolve to the original pages.



**Figure 1.** Deterministic  $C^2$ GES pipeline and path-deletion mechanism. **(a)** Complete reports pass candidate and reference-separation gates; retained passages receive lexical roles and form a typed textual graph. The five score channels feed an ordered role-group reservation followed by redundancy-aware greedy filling. The unrenormalized coefficient-removal comparison sets  $C$  to zero without renormalizing the remaining coefficients. **(b)** For candidate node  $i$ , path-deletion loss removes  $i$ , recomputes qualified-path utility, and takes  $C_i = U(G) - U(G \setminus i)$ .


The unit of analysis is the report. Candidate units within a report share topic, vocabulary, layout, and reference-summary content and are not treated as independent observations. Candidate-level quantities drive selection, whereas paired differences and composition intervals use the 15 report-level observations.

### 3.2. Corpus Construction and Reference Separation

The corpus builder inspected 40 complete public PDFs covering 3200 declared pages. It retained 27 reports and excluded 13 before outcome computation; 12 retained reports were assigned to development and 15 to test. Figure 2 summarizes this construction, and Supplementary Table S1 provides the non-verbatim sampling frame and exclusion reasons.

Inclusion required a complete readable PDF, a detectable summary reference, a non-empty post-summary candidate region, and passage of the predefined extraction-quality gates. Exclusion was decided before outcome computation and remained visible in the inventory. This procedure reduces discretionary removal after observing ROUGE values, but it also changes the estimand: the results concern reports that satisfy the builder's document and reference requirements, not every NERC publication. An operationally important report without a conventional Executive Summary remains outside the present automatic evaluation.

Development and test assignment was performed at report level so that sentences from one source could not cross partitions. Reference-summary text was used only for development scoring or final evaluation, never as a candidate. The builder preserved report ID, page interval, source order, and candidate ID so that every downstream score could be joined directly to its extraction record.



Rights-safe Table S1 records genre, pages, reference words, candidates, split, exclusion reason, and permission status.

**Figure 2.** Corpus construction and partitioning. Boundary and extraction-quality gates retained 27 of 40 complete PDFs. Twelve reports supported configuration selection, and 15 formed the retained test set.

Candidates begin after a conservatively detected Executive Summary boundary, and no fixed sentence cap is applied. The builder produced 12,924 candidates, with 51–1898 per report. Rebuilding the corpus found no reference/body page overlap, normalized exact candidate/reference match, or normalized common substring of at least 50 characters. These checks prevent direct reference leakage caused by page or extraction errors; legitimate paraphrase and facts repeated in the report body remain part of the task.

3.3. Confirmatory Extension Protocol and Layout-Aware Units

The repository also contains a machine-readable protocol for a possible confirmatory extension. It is intentionally retained as a non-executable draft because the reports used in the present pilot were accessed before a freeze and cannot be relabeled as unseen. Any future confirmatory evaluation must instead supply a new rights-cleared inventory, source-file hashes, series-disjoint development and test assignments, immutable method revisions, parameter spaces, random seeds, and a signed tuning decision before outcome access. At least eight independent external series are required and 12 are recommended; report series, rather than individual reports, would be the allocation and inferential unit. This protocol is future-work infrastructure and is not a submission claim gate for the diagnostic study reported here.

The exploratory pilot used a PyMuPDF block rule that distinguishes body paragraphs, headings, list items, table units, captions, and footnotes; removes recurrent headers, footers, and page numbers; prevents table rows from being fused with body sentences; repairs cross-line units; and stores unit type, page, coordinates, source order, word count, and tokenizer length. Three revisions were made after machine error-discovery feedback, so the resulting boundary audit is developmental rather than blinded validation. A future confirmatory study would freeze the rule only after independent boundary review and before external outcomes are accessed.

The exploratory word budgets are 110 and 260 words. Budget enforcement occurs during selection over the complete ranking: a unit that cannot fit the remaining budget is skipped, later units remain eligible, and no unit is truncated. All systems use the same candidates and budget function. Fixed  $K \in \{5, 10\}$  results are retained only for historical comparability.

3.4. Sentence Roles and Typed Textual Proxy Graph

Each candidate receives lexical evidence for root cause, trigger event, propagation or response, impact, and mitigation. Positive top-score ties lead to abstention rather than

resolution by a fixed role order; all 111 development ties received no dominant role and no incident directed edge. The v0.3 selector received no reference role labels.

For each role, the implementation counts predefined cue classes and divides all non-negative scores for that role by the largest score in the report. The role channel  $R_i$  is the maximum of these five unit-scaled values. A dominant role is assigned only when one positive role score is uniquely maximal. An all-zero vector or equal positive maximum yields abstention. An abstaining unit remains eligible for selection through the non-role channels but contributes no typed-path edge.

Table 2 gives the executable taxonomy and representative substring-based lexical cues. The system does not detect grammatical scope, factual status, or document-level coreference.

**Table 2.** Executable role and transition taxonomy. These are lexical text proxies, not expert labels or physical relations.

| Role                    | Representative cue classes                                                 | Permitted outgoing transitions               |
|-------------------------|----------------------------------------------------------------------------|----------------------------------------------|
| Root cause              | cause, due-to, failure, and protection-setting terms                       | trigger; propagation or response; mitigation |
| Trigger event           | fault, outage, initiation, trip, and fire terms                            | propagation or response; impact              |
| Propagation or response | trip, voltage, frequency, islanding, load-shed, reserve, and flow terms    | impact                                       |
| Impact                  | quantities with grid units, customer, load, loss, and unavailability terms | mitigation                                   |
| Mitigation              | recommendation, corrective, preventive, and standard terms                 | none                                         |
| Abstention              | no positive cue, or more than one equal positive maximum                   | none                                         |

A directed edge is allowed only for predefined stage-monotone role transitions whose candidate positions differ by at most 12. The implementation maps root cause, trigger event, propagation or response, impact, and mitigation to stages 0, 1, 2, 3, and 4, respectively. Let  $d_{ij} = |\text{pos}(i) - \text{pos}(j)|$ , let  $J_{ij}$  be token-set Jaccard overlap, and let  $r_i, r_j \in [0, 1]$  be the unit-scaled dominant-role scores. The edge weight is

$$w_{ij} = \text{round}_{12}[0.45 \exp(-d_{ij}/5) + 0.30J_{ij} + 0.25 \min(r_i, r_j)].$$

Thus  $w_{ij} \in [0, 1]$ . Min-max scaled weighted degree within the report forms graph salience  $G_i$ .

The edge gate has semantic and positional components. The source and target roles must appear in the stage-monotone transition set in Table 2, and their candidate indices must lie within the 12-position horizon. The implementation uses absolute positional distance and therefore does *not* require the source-role unit to occur earlier in document order; edge direction follows the role transition rather than chronology. The weight rewards proximity, lexical continuity, and confident role assignments. A high weight therefore means “nearby, lexically connected, and compatible with the predefined role progression”; it does not indicate that a domain expert verified the relation or its temporal order.

These choices make the graph sparse and document specific. Sparse construction limits path enumeration and improves inspectability, while document-specific normalization avoids comparing raw node degrees across reports of very different size. The design also creates known blind spots. A causal statement whose evidence and consequence are separated by more than 12 candidates cannot form an edge; a relation expressed through coreference may have low token overlap; a role-compatible edge may point against docu-

ment chronology; and a valid sentence with two equally strong role cues abstains. These are algorithmic consequences rather than post hoc explanations of individual outcomes.

### 3.5. Typed Path-Deletion Structural Perturbation

Let  $\mathcal{P}(G)$  be the set of simple, stage-monotone paths with two to four edges beginning at a root-cause or trigger node and ending at an impact or mitigation node. For path  $p$ ,

$$\text{strength}(p) = \left( \prod_{e \in p} w_e \right)^{1/|p|} \frac{\max \text{stage}(p) - \min \text{stage}(p)}{4}. \quad (1)$$

The graph utility and raw deletion score are

$$U(G) = \sum_{p \in \mathcal{P}(G)} \text{strength}(p), \quad C_i = U(G) - U(G_{-i}) = \sum_{p \ni i} \text{strength}(p). \quad (2)$$

Unlike weighted degree,  $C_i$  depends on qualified multi-edge path membership, role order, and the geometric mean of edge weights. A synthetic unit-test counterexample gives two nodes equal weighted degree but unequal qualified-path membership. Unit tests verify deletion-loss equality, deterministic scaling, cache equivalence, and fail-closed work limits. These checks establish software and mathematical non-identity, not physical causality or effectiveness.

The functional has three design properties under a fixed qualified-path set: non-negativity when edge weights are non-negative; nullity for a node that belongs to no qualified path; and additivity over the qualified paths that contain the deleted node. More precisely, if deleting node  $i$  removes exactly those paths containing  $i$  and does not change the weights of paths that remain, then  $C_i = \sum_{p \ni i} \text{strength}(p)$ . This proposition is an accounting identity for a fixed textual graph, not a causal-identification theorem. It can fail as a model of semantic influence when deletion would change role assignments, coreference, discourse interpretation, or the weights of surviving relations; the implemented perturbation intentionally holds those quantities fixed.

Table 3 provides a synthetic walk-through that is independent of the experimental corpus. Consider five source-ordered nodes with roles root cause, trigger, propagation, impact, and mitigation and four compatible edge weights. The four qualified paths that begin at the root-cause or trigger node and end at impact or mitigation have strengths 0.566964, 0.741559, 0.367423, and 0.542282 under Equation (1), giving  $U(G) = 2.218228 \approx 2.218$ . Deleting the propagation node removes all four paths, so  $C_q \approx 2.218$ ; deleting the root-cause node removes only the two paths that start there, so  $C_r = 1.308523 \approx 1.309$ . The example shows why path participation and stage span matter beyond local edge totals.

**Table 3.** Synthetic typed-path example used only to explain the implemented computation; it is not an experimental observation.

| Nodes on path                                               | Edge weights       | Stage-span factor | Path strength |
|-------------------------------------------------------------|--------------------|-------------------|---------------|
| $r \rightarrow t \rightarrow q \rightarrow i$               | 0.8, 0.6, 0.9      | 3/4               | 0.567         |
| $r \rightarrow t \rightarrow q \rightarrow i \rightarrow m$ | 0.8, 0.6, 0.9, 0.7 | 4/4               | 0.741         |
| $t \rightarrow q \rightarrow i$                             | 0.6, 0.9           | 2/4               | 0.367         |
| $t \rightarrow q \rightarrow i \rightarrow m$               | 0.6, 0.9, 0.7      | 3/4               | 0.542         |

The equality  $C_i = \sum_{p \ni i} \text{strength}(p)$  follows because node deletion removes exactly the qualified paths containing  $i$  and leaves the remaining path strengths unchanged. The implementation recomputes or cache-checks the deleted graph, and direct deletion loss,

membership-sum loss, and cached loss must agree within the numeric tolerance. The raw  $C_i$  values are then min-max scaled within the report.

### 3.6. Scoring and Selection

The Full base score is

$$S_i = 0.40Q_i + 0.20R_i + 0.15G_i + 0.15C_i + 0.10P_i, \quad (3)$$

where  $Q_i$  is min-max scaled centroid-style lexical relevance,  $R_i$  is the maximum unit-scaled role evidence,  $G_i$  is min-max scaled weighted-degree salience,  $C_i$  is min-max scaled path-deletion loss, and  $P_i = 1/(1 + \text{pos}(i))$  is the unrescaled positional prior for zero-based source position. Path and expansion caps are 250,000 and 2,000,000; reaching either cap terminates the computation rather than returning a partial path score.

The unrenormalized coefficient-removal comparison changes only the coefficient of  $C_i$  from 0.15 to zero. It does not renormalize any remaining coefficient or alter candidates, graph construction, role-group reservation, redundancy penalty, tie rules, or budgets. Consequently, the positive-channel weight decreases from 1.00 to 0.85 while the redundancy coefficient remains 0.50; relative to the remaining positive scale, that penalty is  $0.50/0.85 = 0.5882$ , 17.6% larger than 0.50. The primary contrast therefore estimates coefficient removal under this unrenormalized scoring rule, including relative-scale coupling, rather than a pure information-only contribution of  $C_i$ . A post-run clean diagnostic divides the corresponding no-path base score by 0.85, giving unit-sum remaining weights (0.470588, 0.235294, 0.176471, 0, 0.117647) while retaining the 0.50 redundancy penalty and every selection rule.

Each score channel has a distinct interpretation.  $Q_i$  rewards lexical terms that recur across the report and can therefore favor generic report language.  $R_i$  rewards predefined condition, event, impact, and mitigation cues but is vulnerable to cue ambiguity.  $G_i$  rewards local connectivity under the typed edge rules, whereas  $C_i$  rewards membership in qualified paths and can differ from degree.  $P_i$  favors earlier source positions. The scaling rules bound all channels to  $[0, 1]$  within a report but do not make their meanings interchangeable.

Selection has two phases. When  $K \geq 3$ , the selector first considers three role groups in the fixed order (i) root cause or trigger event, (ii) propagation or response or impact, and (iii) mitigation. If a group has eligible candidates, it reserves the candidate with the largest base score. Ties use earlier source position and then sentence identifier. A group with no eligible candidate is skipped; the method never manufactures a role assignment to satisfy the reservation.

After reservation, the selector fills the remaining positions by maximizing

$$S'_i = S_i - 0.50 \max_{j \in A} J(i, j), \quad (4)$$

where  $A$  is the selected set and an empty maximum is zero. Exact ties again use earlier source position and then sentence identifier. The selected set is finally reordered by source position for presentation. The role-group phase uses the base score without a redundancy penalty, so its contribution to structural coverage is separate from the path-deletion channel. No stochastic decoding or model sampling occurs in this selector.

The dominant computational cost arises from candidate-pair edge checks and bounded path enumeration. The implementation screens every ordered source-target pair, so edge construction is  $O(n^2)$ . The absolute 12-position gate limits admitted local ordered pairs to at most approximately  $24n$  apart from boundary effects, but it is applied after the quadratic scan. Path enumeration can still grow combinatorially in a dense compatible region, which motivates the 250,000-path and 2,000,000-expansion limits.

The output is designed for source navigation. Each selected item preserves report and candidate identifiers, and a deterministic non-verbatim locator joins them to page, source order, report metadata, and source URL. All 1575 selected references resolved uniquely to an integer page within the declared PDF range, allowing readers to reopen the source page and inspect the surrounding context.

### 3.7. Ablation Design and Controlled Variants

The exploratory component study implements three controlled comparison families whose configuration registry is preserved with the results. The incremental chain begins with relevance and position (AB-0), then adds role score (AB-1), role-group reservation (AB-2), typed-graph salience (AB-3), normalized path deletion without redundancy (AB-4), and redundancy to form normalized Full (AB-5). AB-6 retains every AB-5 component except path deletion and is therefore the primary controlled path comparison. The historical unrenormalized coefficient-removal condition receives no AB identifier and remains a scale-coupled historical diagnostic.

A  $2 \times 2$  factorial holds typed graph salience, redundancy, candidates, budgets, seeds, and tie rules fixed while independently switching role-group reservation and normalized path deletion: RP-00 disables both, RP-10 enables reservation only, RP-01 enables path deletion only, and RP-11 enables both. The analysis estimates the reservation main effect, path main effect, and their interaction. A separate G-U/G-T pair holds reservation on, path off, and redundancy on while changing only untyped versus typed graph salience. This pair isolates the net value of edge-type information from graph connectivity itself.

For every clean deletion or addition of a positive component, positive weights are renormalized to sum to one and the redundancy penalty retains the same definition relative to that positive scale. Every condition stores per-report selections, metrics, realized words, budget utilization, runtime, peak memory, failure state, configuration, seed, and code revision. Required outcomes are ROUGE-L, redundancy, role coverage, typed-edge coverage, and selection Jaccard; ROUGE alone is insufficient to interpret a structural component.

### 3.8. Method Assumptions and Failure Modes

The lexical roles are executable features rather than expert labels. Ambiguous positive ties lead to abstention, but an unambiguous cue can still assign an incorrect role. Negation scope, hypothetical or quoted events, and document-level coreference are not modeled. A transition can also join locally compatible roles without a supported discourse relation. These assumptions make the method deterministic and inspectable while limiting semantic coverage.

Errors can enter through segmentation, role assignment, edge construction, or final selection. For example, a PDF table may remain fused with prose, a candidate may receive the wrong dominant role, compatible role cues may form an unsupported edge, or a relevant extract may omit necessary context. ROUGE cannot distinguish these layers. Future expert annotation should therefore score extraction-unit validity, role or abstention, relation support, source faithfulness, and omission separately.

### 3.9. Human-Validation Protocol Reserved for Future Work

No human annotation was conducted for the reported study. The repository preserves a future-work protocol requiring two independent annotators, including at least one person with documented power-system reliability, event-analysis, or closely related technical-report experience. Its target sample comprises 200 extraction units for boundary assessment, 200 role judgments that may reuse those sampled units as a separate task, 150 typed edges, 80 complete typed paths, and 30 system-summary sets, stratified by report series, split, layout, role confidence, path availability, direction agreement, and Full/no-path selection

agreement. During independent annotation, annotators would be blinded to system scores, condition names, automated role labels, and one another's labels. Recruitment and data collection require a documented institutional ethics decision before they begin.

The reserved tasks separately assess standalone unit validity; one of five roles, none/other, or ambiguous/multiple roles; edge support and direction; path coherence and information added beyond role reservation; and source faithfulness, contextual distortion, critical omission, and page-locator correctness. Until that study is independently executed, lexical roles, typed edges, and typed paths are described only as heuristic textual proxies. The two-model audit reported below is kept outside this protocol and cannot satisfy any human-validation threshold.

### 3.10. Development Selection and Comparators

The 12-report development search evaluated 144 predefined configurations. Its ordered objective maximized mean ROUGE-L F1 at  $K = 5$ , followed by mean ROUGE-1 F1, lower redundancy, and earliest configuration order. Grid index 60 was selected, with mean development ROUGE-L F1 of 0.1195126 and ROUGE-1 F1 of 0.2713870. Full minus the unrenormalized coefficient-removal comparison was  $-0.0056652$  on development.

Seven conditions share the same candidate space: Lead, lexical Centroid, TextRank, Semantic-MMR, Role, the strict graph variant without path deletion, and Full C<sup>2</sup>GES. Semantic-MMR uses the local sentence-transformers/all-MiniLM-L6-v2 snapshot at revision 1110a243fdf4706b3f48f1d95db1a4f5529b4d41, with its WordPiece tokenizer, normalized embeddings, CPU execution, and Sentence-Transformers maximum sequence length of 256 tokens. It greedily maximizes the maximal-marginal-relevance objective  $0.5 \cos(e_i, \bar{e}_D) - 0.5 \max_{j \in A} \cos(e_i, e_j)$  [36], evaluated over the Sentence-Transformers representations [5]. Its symmetric coefficient was not optimized on development or test outcomes.

The comparators expose distinct ranking signals. Lead tests position; Centroid tests lexical representativeness; TextRank tests an untyped similarity graph; Semantic-MMR combines dense relevance and diversity; and Role tests lexical discourse evidence without graph structure. The unrenormalized coefficient-removal comparison retains the typed graph while setting the deletion coefficient to zero without renormalization. It includes the scale coupling described in Section 3.5 rather than isolating path information alone.

Hyperparameter opportunity was not balanced for the retained-test comparison (Table 4). Full C<sup>2</sup>GES inherited a 144-configuration development search, whereas the evaluated Semantic-MMR coefficient was fixed at 0.5 and TextRank used its implementation defaults. This asymmetry is relevant to system-level ranks. RQ1 is therefore interpreted as a descriptive comparison of the reported configurations within this corpus, not as a balanced-tuning comparison of algorithms. The unrenormalized coefficient-removal comparison is the closer comparator for RQ2 because it shares the selected absolute weights, graph, candidates, and selector.

**Table 4.** Tuning opportunity and inferential role. The reported intervals condition on the selected Full configuration and do not include model-selection uncertainty.

| Comparison object                  | Development opportunity             | Analytical role                                              |
|------------------------------------|-------------------------------------|--------------------------------------------------------------|
| Full C <sup>2</sup> GES            | 144 predefined configurations       | Selected deterministic system; selection uncertainty omitted |
| Unrenormalized coefficient removal | Same absolute weights, $C = 0$ only | Component comparator; redundancy is relatively stronger      |
| Semantic-MMR                       | Coefficient fixed at 0.5            | Semantic relevance–diversity comparator                      |
| TextRank                           | Implementation defaults             | Untyped graph-ranking comparator                             |

After the retained-test results were known, a separate development-only exercise assigned exactly nine configurations to each of Semantic-MMR, TextRank, and a normalized-path C<sup>2</sup>GES variant under one ordered objective over both budgets. It selected Semantic-MMR coefficient 0.9, TextRank damping 0.65, and normalized path weight 0.0. The script records that it did not access the test input. These choices are reserved for a prospectively frozen external-series study and were not substituted into any retained-test result; consequently, they close the configuration-budget asymmetry only for that planned experiment, not retrospectively for RQ1.

The prospective comparator registry additionally allocates nine configurations to PacSum-MiniLM: preceding-position coefficients  $-2$ ,  $-1$ , or  $0$ , a fixed following-position coefficient of  $1$ , and similarity-threshold fractions  $0$ ,  $0.3$ , or  $0.6$  [14]. Its clean-room source and identity record are frozen in the reproducibility package; the inspected public upstream repository supplied no license file and no upstream code was copied. PacSum-MiniLM tuning and performance evaluation remain blocked until the layout boundary audit passes, so it does not appear in the historical Results.

Sentence-count budgets allow every method to return five or ten complete source units. They do not guarantee equal information because extraction units vary in length. Output length was therefore analyzed after selection as a diagnostic; the frozen primary outputs were neither truncated nor reselected. A separate post-run word-cap sensitivity, defined after the primary result was visible, is described below and remains outside the confirmatory family.

3.11. Statistical Analysis

For the reported post-access pilot, the primary exploratory estimand is the equal-series mean paired ROUGE-L F1 difference. Report outcomes are first averaged within report\_series\_id; series then receive equal weight. The system family compares no-path C<sup>2</sup>GES with Semantic-MMR, TextRank, and PacSum-MiniLM at 110 and 260 words. The component analysis treats the AB incremental chain, reservation–path factorial, and typed–untyped graph comparison as separate families. Every key contrast reports the paired effect, a 95% series-cluster bootstrap interval, an exact series sign-flip value, Holm adjustment within its family, and leave-one-series-out stability. These families were fixed for the v3 diagnostic run but not before source access, so their results remain exploratory. A future confirmatory study would additionally require a positive mean, an interval lower bound above zero, an adjusted exact value below 0.05, matched realized length, and no leave-one-series-out reversal before making an advantage claim.

The historical retained-test estimand is the unweighted mean report-level ROUGE-L difference on the 15-report set. Three Full-minus-baseline contrasts were defined at each budget: the unrenormalized coefficient-removal comparison, Semantic-MMR, and

TextRank. Paired percentile composition intervals use 10,000 report-level resamples. For paired deltas  $d_1, \dots, d_{15}$ , each resample draws 15 reports with replacement and yields mean  $\bar{d}_b$ . The archived analysis also records

$$t_{\text{boot}} = 2 \min \left\{ B^{-1} \sum_{b=1}^B \mathbb{I}(\bar{d}_b \leq 0), B^{-1} \sum_{b=1}^B \mathbb{I}(\bar{d}_b \geq 0) \right\}, \quad B = 10,000. \quad (5)$$

Because resamples are centered on the observed deltas rather than generated under a null,  $t_{\text{boot}}$  is a descriptive sign-tail quantity and not a calibrated hypothesis-test  $p$  value. The composition intervals characterize sensitivity to the retained report set; they are not population confidence intervals. ROUGE-1, ROUGE-2, redundancy, and other pairings are descriptive [37].

An exact paired sign-flip sensitivity enumerates all  $2^{15} = 32,768$  report-level sign assignments for each contrast and applies Holm adjustment across the six values [38]. A second post-result sensitivity maps the 15 reports to the 10 frozen `report_series_id` clusters, averages within each series, assigns equal weight to series, enumerates all  $2^{10} = 1024$  series-level sign assignments, and uses a 10,000-draw series-cluster bootstrap plus leave-one-series-out analysis. A third post-run sensitivity derives 110- and 260-word caps mechanically from development candidate lengths and greedily admits complete sentences from each archived top-10 ranking; its six Full-minus-baseline contrasts use the same equal-series bootstrap, exact series sign flip, and Holm family. None of these sensitivities is a fresh confirmatory test. The interpretation in the main text relies primarily on effect sizes and explicitly distinguishes report-level from series-level robustness.

The run contains all  $15 \times 7 \times 2 = 210$  report-condition-budget records. Packaged checks rebuild aggregate values, contrast records, candidate boundaries, and selected source identifiers. Detailed file identities and check results are provided in Supplementary Table S3.

### 3.12. Exploratory Development-Only Calibration

After the test result was known, a separate exploratory program evaluated 147 semantically deduplicated configurations using only the 12 development reports. It varied path-deletion weight, weight-allocation rule, path horizon, and gating rule. Twelve leave-one-report-out folds selected a zero-weight configuration in 12/12 folds. The best nonzero candidate used weight 0.025, won no fold, and remained below strict zero at both budgets. This analysis diagnoses integration behavior and was not applied to the retained test reports.

### 3.13. Held-Out Matched-Budget Path-Utility Revision

A later freeze-before-eval protocol redesigned the path channel without reading 15-test or seven-series scores for the new utilities. Five utilities (historical deletion loss, relevance-gated, end-stage weighted, length-damped, and gated-length) and three path weights (0.05, 0.10, 0.15) were scored on the 12-report development split at complete-ranking 110- and 260-word budgets. The protocol maximized the development mean of the two budgets, with documented tie-breakers. The frozen choice was the historical utility at weight 0.10. That configuration, no-path C<sup>2</sup>GES, and TextRank were then scored on the 15-report retained test at the same budgets. The 15-test run is held-out for this revision. It is not an unseen confirmatory corpus: those reports were already used in the historical evaluation, and the seen-exclusion registry contained no unused lawful public series.

3.14. *Generative-AI Assistance and Verification Boundary*

The OpenAI Codex agent interface assisted with prose drafting and editing, suggestions for experiment and audit planning, Python and  $\text{\LaTeX}$  development and review, provenance records, manuscript builds, and visual-quality checks. DeepSeek and Codex also independently labeled a small, fixed unit-validity and role packet to identify extraction failures during the exploratory pilot. These labels are reported only as a machine error-discovery audit: neither model was treated as a qualified power-grid expert, human annotator, or adjudicator. All ROUGE, budget, ablation, and statistical values came from documented local execution rather than generated model assertions. The authors reviewed the assisted work and retain responsibility for the study and manuscript.

4. **Results**

4.1. *Overview of Historical Empirical Findings*

The evidence is presented in three explicitly separated layers. The retained-test layer diagnoses the original evaluation, where the unrenormalized coefficient-removal comparison produced the highest overall ROUGE-L mean and Full  $C^2\text{GES}$  exceeded Semantic-MMR and TextRank only under unequal information budgets. The post-access exploratory layer tests the revised word-budget pipeline on seven additional series and supplies controlled component diagnostics. A third layer freezes a redesigned path utility on the 12-report development split and scores the 15-report test at matched 110- and 260-word budgets; that layer is held-out for this revision, not unseen confirmatory. None of the three layers supports confirmatory superiority or semantic-validity claims.

4.2. *Post-Access Exploratory External Pilot*

Eight official reports were fixed in an exploratory inventory after source metadata or snippets had already been observed. Seven PDFs from ENTSO-E, NERC, and Transpower were acquired; the fixed NESO item returned HTTP 403 and was recorded as an acquisition failure rather than replaced. Official summary sections were separated automatically as references without human adjudication. Three rule revisions used machine unit-validity feedback to split heading-body and caption-body fusions, remove incomplete and duplicate body units, and exclude raw table rows from the main experiment. The final dataset contained seven reports in seven independent series and 58–1408 candidates per report. Because access preceded freezing and extraction rules were revised during the exploratory cycle, none of these outcomes is confirmatory.

Table 5 reports the final v3 comparison. All methods used the same candidates and complete-ranking word-budget selector; mean realized length was within 5% of its nominal budget for every method. TextRank ranked first at 110 words, whereas no-path  $C^2\text{GES}$  ranked first at 260 words. For no-path  $C^2\text{GES}$  versus Semantic-MMR, TextRank, and PacSum-MiniLM, all six series-cluster intervals crossed zero and all Holm-adjusted exact values were 1.0. Thus the pilot does not establish a length-controlled advantage.

**Table 5.** Exploratory external-series mean ROUGE-L F1 under matched word budgets ( $n = 7$  series). The highest observed value per budget is bold, but no predefined system contrast survived Holm adjustment.

| Method                     | 110 words     | 260 words     |
|----------------------------|---------------|---------------|
| Lead                       | 0.1657        | 0.1643        |
| Centroid                   | 0.1560        | 0.1678        |
| TextRank                   | <b>0.1690</b> | 0.1796        |
| Semantic-MMR               | 0.1590        | 0.1751        |
| Role-only                  | 0.1563        | 0.1687        |
| No-path C <sup>2</sup> GES | 0.1668        | <b>0.1844</b> |
| Full C <sup>2</sup> GES    | 0.1588        | 0.1772        |
| PacSum-MiniLM              | 0.1661        | 0.1787        |

Full-minus-no-path ROUGE-L was  $-0.00799$  at 110 words (95% series-cluster interval  $[-0.01949, 0.00149]$ ) and  $-0.00718$  at 260 words ( $[-0.01828, 0.00147]$ ); neither exact comparison survived Holm adjustment. Typed-minus-untyped graph effects were  $+0.00139$  and  $-0.00544$ , with both intervals crossing zero. A final machine-only audit of 29 selected units produced 93.1% unit-validity agreement ( $\kappa = 0.633$ ) with no malformed label from either model. Role agreement between the models was 75.9% ( $\kappa = 0.698$ ), but their agreement with the lexical role heuristic was only 48.3% and 65.5%. The layout filter therefore passed this limited machine prescreen, whereas role construct validity remains unestablished.

4.3. Exploratory Component Factorial Ablation

Table 6 reports the controlled path, reservation-path, and graph-type conditions. AB-5 and RP-11 are the same normalized Full configuration; AB-6, RP-10, and G-T are the same typed no-path configuration. They are repeated to make each comparison family explicit. All conditions used identical candidates, budgets, tie rules, and evaluation code, and every one of the 182 report-condition-budget cells completed successfully.

**Table 6.** Exploratory component-factorial results over seven post-access report series. R-L is ROUGE-L F1, Red. is redundancy, Role is role-group coverage, and Edge is typed-edge coverage. Coverage measures formal heuristic structure, not human-validated semantic correctness.

| Condition      | 110-word budget |        |       |       | 260-word budget |        |       |       |
|----------------|-----------------|--------|-------|-------|-----------------|--------|-------|-------|
|                | R-L             | Red.   | Role  | Edge  | R-L             | Red.   | Role  | Edge  |
| AB-5 (Full)    | 0.1588          | 0.0508 | 1.000 | 0.964 | 0.1772          | 0.0509 | 1.000 | 1.000 |
| AB-6 (no path) | 0.1668          | 0.0595 | 1.000 | 0.948 | 0.1844          | 0.0747 | 1.000 | 0.984 |
| RP-00          | 0.1710          | 0.0782 | 0.667 | 0.923 | 0.1870          | 0.0781 | 0.714 | 0.982 |
| RP-10          | 0.1668          | 0.0595 | 1.000 | 0.948 | 0.1844          | 0.0747 | 1.000 | 0.984 |
| RP-01          | 0.1741          | 0.0571 | 0.714 | 0.931 | 0.1735          | 0.0577 | 0.762 | 0.984 |
| RP-11          | 0.1588          | 0.0508 | 1.000 | 0.964 | 0.1772          | 0.0509 | 1.000 | 1.000 |
| G-U            | 0.1654          | 0.0673 | 1.000 | 0.876 | 0.1898          | 0.0814 | 1.000 | 0.944 |
| G-T            | 0.1668          | 0.0595 | 1.000 | 0.948 | 0.1844          | 0.0747 | 1.000 | 0.984 |

The path term materially changed the selected evidence: mean AB-5/AB-6 selection Jaccard was 0.233 at 110 words and 0.479 at 260 words. Nevertheless, AB-5 had lower mean ROUGE-L at both budgets, and this negative direction persisted after omitting each series in turn; the 110-word LOSO estimates ranged from  $-0.01062$  to  $-0.00329$  and the 260-word estimates from  $-0.00912$  to  $-0.00238$ . Thus the null-adjusted uncertainty does not arise because Full and no-path usually selected the same units.

In the factorial family, reservation increased formal role coverage to 1.0 by construction, whereas enabling path without reservation raised coverage only from 0.667 to 0.714 at

110 words and from 0.714 to 0.762 at 260 words. The path effect was positive without reservation but negative with reservation at 110 words; at 260 words it was negative in both strata. The resulting interaction was negative at 110 words and positive at 260 words, and neither interaction survived Holm correction. These budget-dependent directions do not establish either pure path failure or stable functional redundancy with reservation. Supplementary Figure S4 visualizes the interaction, and Supplementary Table S8 reports the complete AB-0–AB-6 chain.

Typed graph information increased formal edge coverage relative to G-U but did not show a stable ROUGE-L benefit. Mean G-T/G-U selection Jaccard was 0.221 and 0.439, confirming that edge typing changed the selections rather than acting as an inert label. The 110-word graph contrast reversed direction in two of seven LOSO analyses, whereas the negative 260-word direction was preserved in all seven. Across the 26 condition–budget groups, failure rate was zero; mean estimated runtime ranged from 0.2466 to 0.4469 seconds per report and mean Python peak memory from 1.03 to 1.49 MB, excluding the separately loaded embedding model. Full and no-path had nearly identical measured resource use, so the present implementation supplies no efficiency argument for retaining the path term.

4.4. Held-Out Matched-Budget Path Revision

Table 7 reports the frozen redesigned-path comparison on the 15-report retained test at complete-ranking 110- and 260-word budgets ( $n = 15$ ). The frozen utility is historical path-deletion loss at weight 0.10, selected on development before these scores were written. The redesigned path did not meet the pre-specified bounded win of a mean at least as large as no-path C<sup>2</sup>GES at both budgets. It was 0.0002 below no-path at 110 words (0.0652 versus 0.0654) and 0.0023 above at 260 words (0.1044 versus 0.1020). TextRank was highest at 110 words (0.0675). These complete-ranking word-budget means are a different estimand from the historical equal-unit  $K = 5/10$  means (Full 0.1060/0.1276; unrenormalized removal 0.1094/0.1310) and do not replace the seven-series Full-minus-no-path diagnosis (−0.00799/−0.00718; no-path 0.1844 at 260 words).

**Table 7.** Held-out-for-this-revision mean ROUGE-L F1 on  $n = 15$  retained-test reports at matched word budgets. The redesigned path was frozen on the 12-report development split. This table is not an unseen confirmatory evaluation.

| Method                             | 110 words | 260 words |
|------------------------------------|-----------|-----------|
| No-path C <sup>2</sup> GES         | 0.0654    | 0.1020    |
| Redesigned path (historical, 0.10) | 0.0652    | 0.1044    |
| TextRank                           | 0.0675    | 0.1031    |

4.5. Synthetic Stress-Test Factorial

A four-series fictional stress set (8 reports; unused themes relative to the earlier network-calibrated fixtures) was generated locally, passed the frozen distribution gates, and was scored with the same complete-ranking 110/260-word factorial used for the post-access pilot. The extractive references were assembled from the same typed role chain that the selector can recover, so this layer tests software and endpoint alignment rather than unseen real-report validity. Table 8 reports the means. Full C<sup>2</sup>GES exceeded no-path at both budgets (0.3886 versus 0.3750 at 110 words; 0.3481 versus 0.3269 at 260 words). The series-equal Full-minus-no-path means were +0.0136 and +0.0212, and both Holm-adjusted exact values were 1.0. Typed-minus-untyped graph means were positive, with Holm-adjusted exact values of 0.25. These fixtures are not NERC or ENTSO-E reports, are not a human-annotated gold set, and do not replace the seven-series Full-minus-no-path diagnosis (−0.00799/−0.00718) or the 15-report held-out mixed result.

**Table 8.** Synthetic-stress mean ROUGE-L F1 on  $n = 8$  fictional reports (4 series) at matched word budgets. References are extractive reconstructions of the planted role chain. This table is not confirmatory and is not real incident data.

| Method                     | 110 words | 260 words |
|----------------------------|-----------|-----------|
| Full C <sup>2</sup> GES    | 0.3886    | 0.3481    |
| No-path C <sup>2</sup> GES | 0.3750    | 0.3269    |
| G-T (typed, no path)       | 0.3750    | 0.3269    |
| G-U (untyped, no path)     | 0.3002    | 0.2346    |

#### 4.6. Historical Retained-Test Main Comparison

Table 9 contains all seven conditions and four recorded metrics at both budgets. Full C<sup>2</sup>GES exceeded Semantic-MMR and TextRank on mean ROUGE-L at each equal-unit budget but was below the unrenormalized coefficient-removal comparison. Because extraction-unit lengths differ materially, these comparisons are not controlled by word or token count. Figure 3 visualizes the same ROUGE-L means.


**Table 9.** Macro-mean results over  $n = 15$  test reports at equal extraction-unit counts, not equal word counts. “Unrenorm. removal” denotes the unrenormalized coefficient-removal comparison. Bold marks only the observed maximum ROUGE-L within a budget; underlining marks the minimum redundancy. Neither convention is an inferential conclusion.

| $K$ | Condition               | R-1 F1 | R-2 F1 | R-L F1        | Redundancy    |
|-----|-------------------------|--------|--------|---------------|---------------|
| 5   | Lead                    | 0.1633 | 0.0545 | 0.0878        | 0.0668        |
| 5   | Centroid                | 0.1955 | 0.0457 | 0.0939        | 0.2333        |
| 5   | TextRank                | 0.1508 | 0.0356 | 0.0806        | 0.2560        |
| 5   | Semantic-MMR            | 0.1783 | 0.0433 | 0.0853        | <u>0.0237</u> |
| 5   | Role                    | 0.1889 | 0.0557 | 0.0934        | 0.0862        |
| 5   | Unrenorm. removal       | 0.2497 | 0.0648 | <b>0.1094</b> | 0.0798        |
| 5   | Full C <sup>2</sup> GES | 0.2359 | 0.0611 | 0.1060        | 0.0745        |
| 10  | Lead                    | 0.2617 | 0.0811 | 0.1158        | 0.0554        |
| 10  | Centroid                | 0.2809 | 0.0683 | 0.1196        | 0.2011        |
| 10  | TextRank                | 0.2470 | 0.0615 | 0.1156        | 0.2149        |
| 10  | Semantic-MMR            | 0.2840 | 0.0728 | 0.1133        | <u>0.0269</u> |
| 10  | Role                    | 0.3132 | 0.0848 | 0.1257        | 0.0550        |
| 10  | Unrenorm. removal       | 0.3471 | 0.0930 | <b>0.1310</b> | 0.0746        |
| 10  | Full C <sup>2</sup> GES | 0.3366 | 0.0874 | 0.1276        | 0.0667        |

Table 10 reports the post-result series-cluster sensitivity. Equal-series estimates retained the direction of every report-level contrast. For Full minus the unrenormalized coefficient-removal comparison, the series-equal mean remained negative at both budgets. For the unequal-length comparisons, the Holm-adjusted exact series-level values were 0.011719 for Semantic-MMR and TextRank at  $K = 5$ , but 0.125000 and 0.134766 at  $K = 10$ . Thus the larger-budget lexical contrasts were not supported by this multiplicity-adjusted series-level sensitivity, and none of these comparisons resolves the length mismatch.

**Table 10.** Post-result series-cluster sensitivity for Full-minus-baseline ROUGE-L differences. The 15 reports form 10 frozen series. Intervals are percentile cluster-bootstrap intervals over equal-weighted series; exact values enumerate all 1024 series sign assignments and are Holm-adjusted across six contrasts. This analysis is exploratory.

| <i>K</i> | Baseline          | Series-equal mean | 95% cluster interval   | Exact <i>p</i> | Holm <i>p</i> |
|----------|-------------------|-------------------|------------------------|----------------|---------------|
| 5        | Unrenorm. removal | -0.005744         | [-0.015986, 0.002230]  | 0.316406       | 0.316406      |
| 5        | Semantic-MMR      | 0.021864          | [0.015484, 0.028248]   | 0.001953       | 0.011719      |
| 5        | TextRank          | 0.028557          | [0.019927, 0.038486]   | 0.001953       | 0.011719      |
| 10       | Unrenorm. removal | -0.004773         | [-0.010494, -0.000084] | 0.091797       | 0.183594      |
| 10       | Semantic-MMR      | 0.014205          | [0.003626, 0.023764]   | 0.031250       | 0.125000      |
| 10       | TextRank          | 0.011010          | [0.001883, 0.019764]   | 0.044922       | 0.134766      |



**Figure 3.** Mean Executive-Summary ROUGE-L F1 for the seven conditions ( $n = 15$  reports) at equal sentence counts but unequal word counts. Bars are descriptive macro-means; paired distributions appear in Figure 5.

Increasing the budget from five to ten units raised mean ROUGE-L for every condition, as expected when more source content is admitted. The ordering differed across methods: Role moved from 0.0934 to 0.1257 and became the strongest non-graph comparator at  $K = 10$ , whereas Centroid was the strongest non-graph comparator at  $K = 5$ . Results at only one extract length would therefore incompletely characterize these systems. The larger budget is not necessarily operationally preferable because brevity, coverage, and omission risk may be valued differently.

The unrenormalized coefficient-removal comparison had the largest observed ROUGE-1, ROUGE-2, and ROUGE-L mean at both budgets. Semantic-MMR had the lowest redundancy, consistent with its explicit diversity objective and shorter outputs. Full C<sup>2</sup>GES reduced redundancy relative to the unrenormalized comparison at both budgets while also reducing all three ROUGE means. The two configurations therefore occupy different observed overlap–redundancy operating points; this result does not show that path deletion improves a joint objective.


4.7. Historical Output-Length Diagnostic

A post-selection script counted Unicode whitespace-delimited words and characters in every selected unit. Supplementary Table S2 contains all 14 condition–budget groups and per-report records. Table 11 shows the four conditions used in the main interpretation. The long-unit and Table-marker counts are selection instances, so an item selected by more than one condition or budget is counted in each output.

**Table 11.** Descriptive output-length diagnostics. “> 100” counts selected-unit instances longer than 100 words; “Table” counts units containing the exact case-sensitive string.

| K  | Condition               | Mean words | Mean characters | > 100 | Table | Maximum words |
|----|-------------------------|------------|-----------------|-------|-------|---------------|
| 5  | Full C <sup>2</sup> GES | 287.7      | 1843.6          | 12    | 14    | 140           |
| 5  | Unrenorm. removal       | 329.0      | 2099.9          | 15    | 16    | 270           |
| 5  | Semantic-MMR            | 184.7      | 1242.1          | 5     | 9     | 138           |
| 5  | TextRank                | 177.0      | 1125.1          | 1     | 7     | 124           |
| 10 | Full C <sup>2</sup> GES | 568.9      | 3619.5          | 25    | 26    | 270           |
| 10 | Unrenorm. removal       | 596.9      | 3822.5          | 28    | 26    | 270           |
| 10 | Semantic-MMR            | 354.5      | 2361.2          | 7     | 16    | 138           |
| 10 | TextRank                | 369.0      | 2329.7          | 4     | 18    | 130           |

Figure 4 makes the output-length mismatch visible. At  $K = 5$ , Full used 56% more words than Semantic-MMR and 63% more than TextRank on average; at  $K = 10$ , the corresponding increases were 61% and 54%.

**Figure 4.** Mean selected words per test report for the four principal conditions. Values reproduce Table 11; no word-budget reselection was performed.

Full averaged 103.0 and 214.5 more words than Semantic-MMR at  $K = 5$  and  $K = 10$ , and 110.7 and 199.9 more than TextRank. Among its 225 selected-unit instances, 37 were longer than 100 words, 40 contained “Table”, and the maximum was 270 words. These diagnostics show fused table, heading, footnote, or line-break blocks among the candidate extraction units. The higher historical ROUGE values relative to Semantic-MMR and TextRank are consequently equal-unit but unequal-word comparisons and do not establish a fair length-matched advantage. They motivated the explicitly post-run word-cap sensitivity in Table 12.

A future confirmatory word-budget study must specify the budget before ranking, define how partially fitting units are handled, and apply an equivalent constraint to every comparator. Truncating the present outputs would answer a different question because it could cut a technical statement mid-clause and would preserve any advantage obtained while ranking long fused units.

For the post-run sensitivity, the median word totals of the first five and ten development candidates were 113.5 and 260.0; rounding to the nearest ten fixed caps of 110 and 260 words. Each method’s archived top-10 ranking was traversed once, admitting only complete sentences that fit the remaining cap and then restoring document order. This reduced mean realized-length differences: at 110 words, condition means ranged from 100.5 to 105.6 words; at 260 words, they ranged from 234.1 to 249.9. The analysis nevertheless searches only the frozen top-10 pool and was designed after test outcomes were known.

All six intervals cross zero and no exact comparison survives Holm adjustment. Thus the top-10 word-cap sensitivity supplies no multiplicity-adjusted evidence for a Full ad-

**Table 12.** Post-run matched-word sensitivity using equal-weight report-series estimands. The 15 reports form 10 frozen series. Intervals use 10,000 series bootstrap draws; exact values enumerate all 1024 series sign assignments and are Holm-adjusted across six contrasts.

| Word cap | Contrast                 | Equal-series difference | Cluster-bootstrap 95% interval | Holm $p$ |
|----------|--------------------------|-------------------------|--------------------------------|----------|
| 110      | Full – unrenorm. removal | +0.0020                 | [−0.0002, 0.0041]              | 0.796875 |
| 110      | Full – Semantic-MMR      | +0.0019                 | [−0.0034, 0.0079]              | 1.000000 |
| 110      | Full – TextRank          | −0.0008                 | [−0.0066, 0.0046]              | 1.000000 |
| 260      | Full – unrenorm. removal | −0.0015                 | [−0.0072, 0.0021]              | 1.000000 |
| 260      | Full – Semantic-MMR      | +0.0062                 | [−0.0006, 0.0142]              | 0.796875 |
| 260      | Full – TextRank          | +0.0001                 | [−0.0065, 0.0059]              | 1.000000 |

vantage over any of the three principal comparators. Because it was post-run and does not rank beyond the archived top ten, it narrows the length concern but cannot replace a prospectively frozen matched-budget experiment.

The long-unit counts locate a concrete extraction risk. Units exceeding 100 words and units containing “Table” show that PDF layout sometimes survives the sentence splitter as one selectable block. Such blocks can contain many reference-summary tokens and thereby increase ROUGE without providing a concise sentence. They may remain useful source locators, but they weaken sentence-budget comparisons and motivate layout-aware segmentation before external validation.

A nonverbatim post-result extraction audit subsequently applied PyMuPDF block coordinates to all 27 included PDF files, with every source hash matching the frozen ledger. The rule preserved page/block boundaries, excluded recurrent margin material, and removed blocks overlapping detected tables by at least 25%. It detected 505 table regions with zero page-level detection failures and produced 14,290 candidate units, of which 9824 were exact normalized matches to the 12,924 legacy units. The number of candidates exceeding 100 words decreased from 214 in the legacy pool to 39 under the block-preserving rule. This is an extraction-boundary diagnostic, not a performance result: candidates and references were not ranked or scored, and one retained block still reached 522 words. Hashed first, median, and longest-unit audit rows support manual checking without redistributing report text.

A post-run tokenizer audit measured all 12,924 development and test candidates without truncation to quantify exposure to the frozen encoder limit. Median length was 33 tokens, the 95th percentile was 89, and 38 candidates (0.294%) exceeded 256 tokens. On the retained test split, 29 of 9504 candidates (0.305%) exceeded 256 and the maximum was 525. A subsequent representation sensitivity first reproduced every frozen production Semantic-MMR selection and metric, then compared the same model with a 512-token limit and with 254-content-token chunk embeddings averaged only for over-limit candidates. The 512-token variant changed 0/15 reports at  $K = 5$  and 2/15 at  $K = 10$ ; chunk averaging changed 1/15 and 2/15. Equal-series ROUGE-L changes ranged from  $-0.00057$  to  $+0.00058$ , and all four exact series comparisons had Holm-adjusted  $p = 1.0$ . These tested alternatives had limited impact on this retained corpus, but they were defined post-run and do not establish general truncation robustness or identify a preferred long-text representation.

#### 4.8. Historical Endpoint Contribution of the Path-Deletion Term

Table 13 gives the six predefined contrasts. Full minus the unrenormalized coefficient-removal comparison was  $-0.003332$  at  $K = 5$  and  $-0.003360$  at  $K = 10$ ; both report-composition intervals crossed zero. Full had slightly lower redundancy than the unrenormalized comparison (0.0745 versus 0.0798 at  $K = 5$  and 0.0667 versus 0.0746 at  $K = 10$ ), but redundancy was a descriptive secondary outcome. Under the current integration, the path-deletion term therefore did not improve the predefined lexical-overlap endpoint.


Full also exceeded Semantic-MMR and TextRank at both equal-unit budgets on this split. These paired differences describe the selected reports but retain the output-length mismatch documented in Section 4.3. No minimum important difference or qualified-reader outcome was specified, so the automatic contrasts do not establish practical benefit for maintenance review.

The post-run clean normalized ablation reproduced every archived Full and unrenormalized coefficient-removal selection before constructing the new comparator. Normalization changed that selection on 1/15 reports at  $K = 5$  and 3/15 at  $K = 10$ . Full minus normalized no-path was  $-0.00613$  and  $-0.00393$  under equal-series weighting; the series-cluster intervals were  $[-0.01591, 0.00166]$  and  $[-0.01006, 0.00104]$ , and both Holm-adjusted exact values were  $0.441406$ . This same-corpus diagnostic separates path removal from the  $0.85$  positive-score scale but still supplies no multiplicity-adjusted evidence that the path term improves ROUGE-L.

**Table 13.** Predefined Full-minus-baseline Executive-Summary ROUGE-L contrasts paired by test report ( $n = 15$ ). Intervals describe sensitivity of the mean paired difference to the composition of the retained reports. The final column is a descriptive bootstrap sign-tail quantity rather than a calibrated hypothesis-test  $p$  value.

| $K$ | Baseline          | Mean difference | 95% comp. interval    | Unadjusted $t_{boot}$ |
|-----|-------------------|-----------------|-----------------------|-----------------------|
| 5   | Unrenorm. removal | -0.003332       | [-0.010889, 0.002826] | 0.3544                |
| 5   | Semantic-MMR      | 0.020737        | [0.012765, 0.028354]  | 0/10,000 tail draws   |
| 5   | TextRank          | 0.025438        | [0.017542, 0.033722]  | 0/10,000 tail draws   |
| 10  | Unrenorm. removal | -0.003360       | [-0.008306, 0.001040] | 0.1444                |
| 10  | Semantic-MMR      | 0.014360        | [0.006082, 0.022215]  | 0.0008                |
| 10  | TextRank          | 0.012029        | [0.004397, 0.019241]  | 0.0022                |

Figure 5 shows every report-level difference. Full and the unrenormalized coefficient-removal comparison selected different candidate sequences in 28 of 30 report-budget cells, and their base scores differed in all 30. The coefficient rule was therefore active, although activity alone does not indicate a benefit.



**Figure 5.** Per-report paired Executive-Summary ROUGE-L differences for the six predefined contrasts ( $n = 15$ ). Horizontal colored segments denote means; black lines denote zero; panel titles give positive/negative/zero sign counts. Rights-safe report indices map to non-verbatim metadata in Supplementary Table S1.

The report-level signs show heterogeneity hidden by the means. Against the unrenormalized coefficient-removal comparison, Full was positive and negative on equally many non-tied reports at  $K = 5$  and negative on more reports at  $K = 10$ . Against Semantic-MMR and TextRank, positive differences predominated, although the length mismatch remains. Report-level and series-level sign-flip sensitivities produced the same directions under their respective sign-exchangeability assumptions, but the multiplicity-adjusted series-level values did not support the  $K = 10$  baseline contrasts. Full report-level values are in Supplementary Table S4; the series-level machine-readable analysis is included with the reproducibility files.

#### 4.9. Historical Component Behavior and Coefficient Support

Across 19,008 Full/unrenormalized-removal candidate-score comparisons, 9774 were nonzero, and selected sequences changed in 28 of 30 report-budget cells. The term was therefore numerically active. Its report-level effect was heterogeneous, and the endpoint comparison in Section 4.4 did not show an average ROUGE-L gain.

The post-result development-only calibration evaluated 147 semantically deduplicated configurations. A zero path-deletion weight won all 12 leave-one-report-out folds. The best nonzero candidate used weight 0.025, won no fold, and remained below its zero-coefficient comparison at both budgets. The separate equal-nine development exercise likewise selected normalized path weight 0.0 while selecting Semantic-MMR coefficient 0.9 and TextRank damping 0.65 for a future external-series protocol. Neither exercise was applied to the retained test. Together they make a simple coefficient adjustment an unpromising explanation for the negative ablation, but they cannot replace an untouched comparison. A revised component will likely require different path semantics, role evidence, or score integration.

Maximum path expansions were 19,638 of the 2,000,000 cap, and qualifying paths were 16,182 of 250,000. The selected reports were therefore processed without approaching the

work limits. These observed margins describe the present corpus rather than asymptotic scalability.

#### 4.10. Implementation Reliability

All reported selections and aggregate values were reproduced from the archived run artifacts. Candidate boundaries, condition–budget completeness, source identifiers, path-work limits, and aggregate recalculation passed the packaged checks. These tests support the integrity of the reported experiment. The packaged reproducibility records contain the complete matrix, file identities, and corrective history, with key checks summarized in the Supplementary Materials.

## 5. Discussion

### 5.1. Main Findings and System-Level Comparison

The system-level result is mixed. At equal extraction-unit counts, Full C<sup>2</sup>GES recovered more Executive-Summary wording than Semantic-MMR and TextRank, but its outputs contained 54–63% more words. The post-access exploratory pilot corrected that length mismatch and placed no-path C<sup>2</sup>GES close to TextRank and PacSum-MiniLM, but no predefined contrast survived multiplicity adjustment. Full was also below no-path at both exploratory word budgets. Taken together, these results do not establish a system-level advantage and instead favor no-path C<sup>2</sup>GES as the simpler provisional configuration pending a prospectively frozen study on genuinely unseen report series.

### 5.2. Why the Path-Deletion Term Did Not Improve ROUGE-L


The controlled factorial changes the interpretation of the negative component result. The path term was not inert: AB-5 and AB-6 had low-to-moderate selection overlap, and Full remained below no-path in every leave-one-series-out rerun. The problem is therefore not that the implementation silently ignored the feature. Rather, path participation redirected selection toward units that increased formal path or edge coverage without improving Executive-Summary overlap. This distinction matters because a no-effect implementation bug and an active but endpoint-misaligned mechanism require different remedies.

Role-group reservation remains one plausible source of overlap, because it already requires available evidence from condition/trigger, propagation/impact, and mitigation groups before ordinary greedy filling. The 2 × 2 results, however, do not establish stable redundancy: at 110 words the path effect changed from positive without reservation to negative with reservation, whereas at 260 words it was negative in both strata, and neither interaction survived family-wise correction. The defensible conclusion is therefore narrower: reservation and path can interact in a budget-dependent manner, but the seven-series pilot cannot identify a stable substitutive relationship.

Score integration is a second issue. The path channel competes with relevance and graph salience after report-level scaling. In the historical strict removal, deleting its positive weight also left the fixed redundancy penalty relatively stronger; this scale coupling motivated the normalized comparison. Because the normalized result was still negative, rescaling alone does not explain the outcome.

Lexical roles create a second limitation. Negation, hypothetical language, quotations, and cross-unit coreference are not resolved, and an edge is permitted by role compatibility and local distance rather than by expert-validated discourse relations. A candidate can therefore participate in a strong typed path without adding useful Executive-Summary wording. Conversely, a semantically important passage can receive low path participation when its related evidence lies beyond the 12-position horizon. These hypotheses are

compatible with the active but heterogeneous score changes, but targeted evaluation is needed to distinguish among them.



**Figure 6.** Component diagnosis for the path-deletion term. The term altered 9774 of 19,008 sentence scores and 28 of 30 report-budget selections, but Full-minus-strict-variant ROUGE-L was slightly negative at both budgets and zero weight was selected in all 12 development leave-one-report-out folds. Error bars in panel (c) are report-composition intervals.

### 5.3. Methodological Lessons for Structure-Aware Long-Document Extraction

The negative component result has value beyond this implementation only if it yields testable design lessons. First, structural complexity must be evaluated incrementally: a richer graph is not supported merely because it changes rankings or improves an internal coverage score. Its contribution must be isolated against a normalized comparator and judged on a task endpoint or an independently validated structural endpoint. Second, hard coverage constraints and soft structural scores should be crossed factorially. Testing either mechanism only inside a full system can conceal substitution, antagonism, or budget dependence. Third, automatically typed edges require a construct-validity study; formal typed-edge coverage cannot be interpreted as semantic coherence when the types originate from lexical rules. Fourth, matched information budgets are essential in long-document extraction because equal sentence counts can reverse the apparent system ranking by allowing methods to emit different amounts of text.

These lessons define a falsifiable redesign path. A revised path mechanism should use expert-validated relations or uncertainty-aware role probabilities, avoid rewarding unsupported chains, and be frozen before evaluation on untouched series. It should be retained only if it improves either a prospectively declared task endpoint or a human-validated structure endpoint without unacceptable redundancy or cost. Otherwise, the simpler no-path selector is the appropriate default. This turns the failed component from a claim of hidden superiority into a documented boundary condition for future structure-aware summarization work. Stated as an information result rather than an engineering one, the exploratory comparison did not establish that role typing contributes discriminative information beyond the untyped graph: the typed-minus-untyped effects did not agree in sign across the two budgets and both intervals crossed zero. That is a failure to demonstrate added information, not a demonstration of redundancy; no equivalence test was run, and the redesign path below must be read against that weaker claim.

### 5.4. Engineering Use and Source-Linked Navigation

The intended practical role of C<sup>2</sup>GES is to provide a structured, source-linked entry point into long reports. A reader can begin with passages associated with condition, event, impact, and mitigation roles and then return to their original pages. Deterministic selection also makes changes attributable to explicit feature channels rather than generation variability. This combination supports method development and analyst-facing prototypes, although automatic reference overlap is not a substitute for engineering judgment.

NERC reliability and disturbance reports can support lessons-learned review, but work orders and inspection records differ in vocabulary, length, asset granularity, narrative structure, and available reference summaries. An operational study would also need asset-specific terminology, access controls, abstention rules for extraction failures, and explicit evaluation of missed hazards. The current experiment establishes a technical-report workflow, not maintenance effectiveness.

### 5.5. Limitations

The test set contains 15 selected public reports from one organization, and the reports were inspected during corrective reconstruction. A post-result sensitivity grouped them into 10 frozen report-series clusters, but it cannot eliminate recurring institutional language, near-duplicate passages, or dependence between nominally distinct series. Both report- and series-level estimates therefore describe the retained corpus rather than population-wide uncertainty.

Equal sentence counts produced unequal word counts, and some extraction units retained table, heading, footnote, or line-break material. The post-run word-cap analysis equalized caps only within frozen top-10 rankings and cannot remove this extraction-boundary bias. The block-preserving extraction audit reduced long-unit incidence but was not used to rerun ranking or obtain human validity judgments. The 512-token and chunk-mean Semantic-MMR diagnostics changed few selections, but both are same-corpus post-run checks. Semantic-MMR used a fixed coefficient of 0.5 in the retained-test run, whereas C<sup>2</sup>GES inherited a 144-configuration development search, so historical tuning opportunities were asymmetric. The later equal-nine development exercise supplies budget-matched candidate configurations only for a future external-series study. A fairer system comparison still requires prospectively frozen layout-aware units, equal word or token budgets, comparable development resources, and untouched evaluation data.

The method lacks expert role labels and expert assessment of source faithfulness, structural coverage, unsafe omission, and reading utility. ROUGE and redundancy cannot provide these judgments. The exploratory external references were extracted automatically, the source inventory was not unseen when its rules were established, one fixed report could not be acquired, and only seven series were available. The freeze-before-eval path revision uses the same 15 retained-test reports as a held-out split for this revision; it is not an unseen confirmatory corpus. The eight-report synthetic factorial uses fictional fixtures and planted extractive references; it is a software stress check, not real incident evidence. The two-model audit improved boundary-error detection but cannot establish construct validity; its limited agreement with lexical roles reinforces this restriction. Third-party licensing also prevents public redistribution of the source PDFs and verbatim derived text; non-verbatim metadata, configurations, aggregate results, code, and page locators can still be provided for review.

### 5.6. Future Validation

The next evidence cycle is now specified rather than left as a general recommendation. E1 freezes a rights-cleared unseen-series corpus, promotes the audited block/page rule only after boundary review, balances development opportunity, and compares complete-ranking outputs at 110 and 260 words. E3 runs the AB-0–AB-6 chain, the reservation-path 2 × 2 factorial, and the G-U/G-T graph-type comparison on the same candidates and series split. The path component is not redesigned after external outcomes are visible; the protocol instead diagnoses whether its effect is absent, redundant with reservation, or confined to structural coverage.

E2 uses independent, blinded judgments to evaluate unit validity, roles, relations, paths, source faithfulness, and omissions. Agreement is reported before adjudication, and failed construct thresholds force claim downgrading. A license-cleared corpus of maintenance narratives or work orders remains a separate conditional E5 application test. Expert task utility is likewise a separate E4 study and is required before any claim about review time, task accuracy, or unsafe omission. Neither E4 nor E5 is implied by successful completion of E1–E3.

## 6. Conclusions

C<sup>2</sup>GES is a deterministic, role-conditioned framework for extracting source-linked passages from long power-system technical reports. Its diagnostic evaluation yields two clear findings. First, Full C<sup>2</sup>GES recovered more Executive-Summary wording than Semantic-MMR and TextRank at equal unit counts, but it produced longer extracts, and the matched-word sensitivity yielded no Holm-supported contrast. The study therefore does not establish a length-controlled system advantage. Second, the unrenormalized coefficient-removal comparison achieved the highest overall means on the selected 15-report NERC test set, showing that the original ablation was scale-coupled rather than a clean estimate of path information.

The path-deletion term changed candidate scores and selected sequences without improving the predefined retained-test endpoint, and development-only calibration selected zero path weight in all 12 folds. In the post-access exploratory pilot, Full was again below no-path by 0.0080 and 0.0072 ROUGE-L at 110 and 260 words, respectively, while typed graph information showed no stable advantage over an untyped graph. A freeze-before-eval redesign of that path term, chosen on development, then scored 0.0652/0.1044 against no-path 0.0654/0.1020 on the same 15 reports at matched word budgets; it did not win both budgets and is not unseen confirmatory. On fictional stress fixtures whose gold summaries were built from the same role chain, Full exceeded no-path in mean ROUGE-L but Holm-adjusted path contrasts remained 1.0; that alignment does not transfer to real Executive Summaries. The present empirical contribution is therefore an auditable implementation, a portable public verification subset, and a reproducible negative component diagnosis. Read as an information result, the study did not establish that the typed structure adds information beyond the untyped graph; the comparison is a failure to demonstrate an advantage, not a demonstration of equivalence. No-path C<sup>2</sup>GES is the provisional default. System superiority, semantic validity of the structural proxies, and operational benefit are not established by this study.

**Supplementary Materials:** The supplementary PDF contains Tables S1–S8 and Figure S4. Tables S5–S8 document the post-access external-series inventory, comparator controls, principal component contrasts, and complete AB-0–AB-6 chain; Figure S4 shows the reservation–path interaction across ROUGE-L and formal structure-coverage metrics. Machine-readable historical results, selected-ID-to-page locators, selection Jaccard, series effects, leave-one-series-out diagnostics, runtime resources, and explicit human-validation status are included in the accompanying reproducibility files. Machine labels are explicitly separated from human validation. Source PDFs, verbatim derived text, and raw model responses are excluded because third-party redistribution permission has not been established.

**Author Contributions:** Conceptualization, B.L. and Y.Y.; methodology, B.L.; software, B.L.; validation, B.L.; formal analysis, B.L.; investigation, B.L.; resources, Y.Y.; data curation, B.L.; writing—original draft preparation, B.L.; writing—review and editing, B.L. and Y.Y.; visualization, B.L.; supervision, Y.Y.; project administration, Y.Y.; funding acquisition, Y.Y. All authors have read and agreed to the published version of the manuscript.

**Funding:** This research was funded by the Science and Technology Research Project of State Grid Fujian Electric Power Co., Ltd., grant number 521300250006. The funder had no role in the design of

the study; in the collection, analyses, or interpretation of data; in the writing of the manuscript; or in the decision to publish the results.

**Institutional Review Board Statement:** Not applicable. The reported work is a computational analysis of published technical reports and includes no recruited human participants, human annotators, or animals. Any future human-annotation study will require the applicable institutional determination before recruitment or data collection.

**Informed Consent Statement:** Not applicable.

**Data Availability Statement:** The historical protocol-ready package is available in the public powergrid\_benchmark repository under tag c2ges-2026-09-06-protocol-ready-v1 via the [frozen C2GES package](#). The diagnostic-submission release adds the post-access exploratory v3 code, the freeze-before-eval path-utility protocol and held-out matched-budget scores, non-verbatim inventory, layout audits, aggregate results, series-level statistics, selected identifiers, and machine-audit diagnostics. Owing to third-party licensing restrictions, source PDFs, verbatim derived text, and raw model responses are not publicly redistributed. Access to restricted materials is subject to applicable third-party permissions and licenses.

**Acknowledgments:** During preparation of this manuscript, the authors used the OpenAI Codex agent interface for the assisted drafting, coding, and verification activities described in Materials and Methods. DeepSeek v4 Flash and Codex gpt-5.6-luna also supplied independent machine labels for the explicitly non-human exploratory quality audit. The authors reviewed and edited the assisted work and take full responsibility for the content of this publication. Neither system was treated as an author, qualified power-grid expert, human annotator, or adjudicator.

**Conflicts of Interest:** The authors declare no conflicts of interest.

## References

1. Zhong, M.; Liu, P.; Chen, Y.; Wang, D.; Qiu, X.; Huang, X. Extractive Summarization as Text Matching. In Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics. Association for Computational Linguistics, 2020, pp. 6197–6208. <https://doi.org/10.18653/v1/2020.acl-main.552>.
2. Bi, K.; Jha, R.; Croft, W.B.; Celikyilmaz, A. AREDSUM: Adaptive Redundancy-Aware Iterative Sentence Ranking for Extractive Document Summarization. In Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume. Association for Computational Linguistics, 2021, pp. 281–291. <https://doi.org/10.18653/v1/2021.eacl-main.22>.
3. Liu, J.; Hughes, D.J.D.; Yang, Y. Unsupervised Extractive Text Summarization with Distance-Augmented Sentence Graphs. In Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval. Association for Computing Machinery, 2021, pp. 2313–2317. <https://doi.org/10.1145/3404835.3463111>.
4. Mihalcea, R.; Tarau, P. TextRank: Bringing Order into Text. In Proceedings of the 2004 Conference on Empirical Methods in Natural Language Processing, Barcelona, Spain, 2004; pp. 404–411.
5. Reimers, N.; Gurevych, I. Sentence-BERT: Sentence Embeddings Using Siamese BERT-Networks. In Proceedings of the EMNLP-IJCNLP 2019. Association for Computational Linguistics, 2019, pp. 3982–3992. <https://doi.org/10.18653/v1/D19-1410>.
6. Siragusa, G.; Robaldo, L. Sentence Graph Attention for Content-Aware Summarization. *Applied Sciences* **2022**, *12*, 10382. <https://doi.org/10.3390/app122010382>.
7. Onan, A.; Alhumyani, H. Contextual Hypergraph Networks for Enhanced Extractive Summarization: Introducing Multi-Element Contextual Hypergraph Extractive Summarizer (MCHES). *Applied Sciences* **2024**, *14*, 4671. <https://doi.org/10.3390/app14114671>.
8. Hark, C. Using Graph-Based Maximum Independent Sets with Large Language Models for Extractive Text Summarization. *Applied Sciences* **2025**, *15*, 6395. <https://doi.org/10.3390/app15126395>.
9. Liu, P.; Tian, B.; Liu, X.; Gu, S.; Yan, L.; Bullock, L.; Ma, C.; Liu, Y.; Zhang, W. Construction of Power Fault Knowledge Graph Based on Deep Learning. *Applied Sciences* **2022**, *12*, 6993. <https://doi.org/10.3390/app12146993>.
10. Ma, T.; Yu, J.; Wang, B.; Gao, M.; Yang, Z.; Li, Y.; Fan, M. A Power Monitor System Cybersecurity Alarm-Tracing Method Based on Knowledge Graph and GCNN. *Applied Sciences* **2025**, *15*, 8188. <https://doi.org/10.3390/app15158188>.
11. Zhang, Y.; Gao, H. Construction of Bridge Maintenance Knowledge Graph Based on Deep Learning. *Applied Sciences* **2026**, *16*, 1985. <https://doi.org/10.3390/app16041985>.
12. North American Electric Reliability Corporation. Attributes of a Quality Event Analysis Report. Technical report, North American Electric Reliability Corporation, 2015.

13. Erkan, G.; Radev, D.R. LexRank: Graph-based Lexical Centrality as Saliency in Text Summarization. *Journal of Artificial Intelligence Research* **2004**, *22*, 457–479. <https://doi.org/10.1613/jair.1523>. 956
14. Zheng, H.; Lapata, M. Sentence Centrality Revisited for Unsupervised Summarization. In Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics, 2019, pp. 6236–6247. <https://doi.org/10.18653/v1/P19-1628>. 957
15. Koh, Y.; Kang, S.; Lee, S. Deep Learning-Based Bug Report Summarization Using Sentence Significance Factors. *Applied Sciences* **2022**, *12*, 5854. <https://doi.org/10.3390/app12125854>. 958
16. Zhang, X.; Wei, Q.; Song, Q.; Zhang, P. TOMDS (Topic-Oriented Multi-Document Summarization): Enabling Personalized Customization of Multi-Document Summaries. *Applied Sciences* **2024**, *14*, 1880. <https://doi.org/10.3390/app14051880>. 959
17. Wang, D.; Liu, P.; Zheng, Y.; Qiu, X.; Huang, X. Heterogeneous Graph Neural Networks for Extractive Document Summarization. In Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics. Association for Computational Linguistics, 2020, pp. 6209–6219. <https://doi.org/10.18653/v1/2020.acl-main.553>. 960
18. Cui, P.; Hu, L.; Liu, Y. Enhancing Extractive Text Summarization with Topic-Aware Graph Neural Networks. In Proceedings of the 28th International Conference on Computational Linguistics. International Committee on Computational Linguistics, 2020, pp. 5360–5371. <https://doi.org/10.18653/v1/2020.coling-main.468>. 961
19. Jing, B.; You, Z.; Yang, T.; Fan, W.; Tong, H. Multiplex Graph Neural Network for Extractive Text Summarization. In Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing. Association for Computational Linguistics, 2021, pp. 133–139. <https://doi.org/10.18653/v1/2021.emnlp-main.11>. 962
20. Bugueño, M.; Hamdan, H.A.; De Melo, G. GraphLSS: Integrating Lexical, Structural, and Semantic Features for Long Document Extractive Summarization. In Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies, 2025, pp. 797–804. <https://doi.org/10.18653/v1/2025.naacl-short.67>. 963
21. Yuan, H.; Hong, S.; Zhang, H. StrucSum: Graph-Structured Reasoning for Long Document Extractive Summarization with LLMs. In Proceedings of the Findings of the Association for Computational Linguistics: EACL 2026, 2026, pp. 3708–3721. <https://doi.org/10.18653/v1/2026.findings-eacl.192>. 964
22. Feder, A.; Keith, K.A.; Manzoor, E.; Pryzant, R.; Sridhar, D.; Wood-Doughty, Z.; Eisenstein, J.; Grimmer, J.; Reichart, R.; Roberts, M.E.; et al. Causal Inference in Natural Language Processing: Estimation, Prediction, Interpretation and Beyond. *Transactions of the Association for Computational Linguistics* **2022**, *10*, 1138–1158. [https://doi.org/10.1162/tacl\\_a\\_00511](https://doi.org/10.1162/tacl_a_00511). 965
23. Du, L.; Ding, X.; Xiong, K.; Liu, T.; Qin, B. e-CARE: A New Dataset for Exploring Explainable Causal Reasoning. In Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics, Volume 1 (Long Papers). Association for Computational Linguistics, 2022, pp. 432–446. <https://doi.org/10.18653/v1/2022.acl-long.33>. 966
24. Xie, L.; Zheng, X.; Sun, Y.; Huang, T.; Bruton, T. Massively Digitized Power Grid: Opportunities and Challenges of Use-Inspired AI. *Proceedings of the IEEE* **2023**, *111*, 762–787. <https://doi.org/10.1109/JPROC.2022.3175070>. 967
25. Hamann, H.F.; Brunschweiler, T.; Gjorgiev, B.; Martins, L.S.A.; Puech, A.; Varbella, A.; Weiss, J.; Bernabe-Moreno, J.; Masse, A.B.; Choi, S.; et al. Foundation Models for the Electric Power Grid, 2024, [arXiv:cs.LG/2407.09434]. 968
26. Madabhushi, S.; Dewri, R. A Survey of Anomaly Detection Methods for Power Grids. *International Journal of Information Security* **2023**, *22*, 1799–1832. <https://doi.org/10.1007/s10207-023-00720-z>. 969
27. Srinivasan, S.; Kumarasamy, S.; Andreadakis, Z.; Lind, P.G. Artificial Intelligence and Mathematical Models of Power Grids Driven by Renewable Energy Sources: A Survey. *Energies* **2023**, *16*, 5383. <https://doi.org/10.3390/en16145383>. 970
28. Ranawaka, A.S.R.A.; Alahakoon, D.; Sun, Y.; Hewapathirana, K. Leveraging the Synergy of Digital Twins and Artificial Intelligence for Sustainable Power Grids: A Scoping Review. *Energies* **2024**, *17*, 5342. <https://doi.org/10.3390/en17215342>. 971
29. Meng, Q.; Zhang, X.; Dong, Y.; Chen, Y.; Lin, D. A Combined Semantic Dependency and Lexical Embedding RoBERTa Model for Grid Field Relational Extraction. *Applied Sciences* **2023**, *13*, 11074. <https://doi.org/10.3390/app131911074>. 972
30. Li, B.; Wang, W. Knowledge-Graph-Based Integrated Line Loss Evaluation Management System. *Applied Sciences* **2024**, *14*, 9462. <https://doi.org/10.3390/app14209462>. 973
31. Cao, X.; Xu, W.; Zhao, J.; Duan, Y.; Yang, X. Research on Large Language Model for Coal Mine Equipment Maintenance Based on Multi-Source Text. *Applied Sciences* **2024**, *14*, 2946. <https://doi.org/10.3390/app14072946>. 974
32. Di Maggio, L.G. Toward Autonomous LLM-Based AI Agents for Predictive Maintenance: State of the Art, Challenges, and Future Perspectives. *Applied Sciences* **2025**, *15*, 11515. <https://doi.org/10.3390/app152111515>. 975
33. Li, Y.; Wang, H.; Zhang, D. An Entity-Relation Extraction Method Based on the Mixture-of-Experts Model and Dependency Parsing. *Applied Sciences* **2025**, *15*, 2119. <https://doi.org/10.3390/app15042119>. 976
34. She, Y.; Tian, T.; Zhang, J. CLEAR: Cross-Document Link-Enhanced Attention for Relation Extraction with Relation-Aware Context Filtering. *Applied Sciences* **2025**, *15*, 7435. <https://doi.org/10.3390/app15137435>. 977
35. Lin, C.Y. ROUGE: A Package for Automatic Evaluation of Summaries. In Proceedings of the Text Summarization Branches Out, Barcelona, Spain, 2004; pp. 74–81. 978

36. Carbonell, J.; Goldstein, J. The Use of MMR, Diversity-Based Reranking for Reordering Documents and Producing Summaries. In Proceedings of the Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval, 1998, pp. 335–336. <https://doi.org/10.1145/290941.291025>. 1010  
1011  
1012
37. Cameron, A.C.; Gelbach, J.B.; Miller, D.L. Bootstrap-Based Improvements for Inference with Clustered Errors. *The Review of Economics and Statistics* **2008**, *90*, 414–427. <https://doi.org/10.1162/rest.90.3.414>. 1013  
1014
38. Holm, S. A Simple Sequentially Rejective Multiple Test Procedure. *Scandinavian Journal of Statistics* **1979**, *6*, 65–70. <https://doi.org/10.2307/4615733>. 1015  
1016

**Disclaimer/Publisher’s Note:** The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content. 1017  
1018  
1019